

Markets and Trading

FRE6871 & FRE7241, Fall 2025

Jerzy Pawlowski jp3900@nyu.edu

NYU Tandon School of Engineering

December 21, 2025



NYU

**TANDON SCHOOL
OF ENGINEERING**

Downloading Treasury Bond Rates from FRED

The constant maturity Treasury yields are the yields of hypothetical fixed-maturity bonds, interpolated from the market yields of actual Treasury bonds.

The *FRED* database provides historical constant maturity Treasury yields:

<https://fred.stlouisfed.org/series/DGS5>

`quantmod::getSymbols()` creates objects in the specified *environment* from the input strings (names).

It then assigns the data to those objects, without returning them as a function value, as a *side effect*.

```
> # Symbols for constant maturity Treasury rates
> symbolv <- c("DGS1", "DGS2", "DGS5", "DGS10", "DGS20", "DGS30")
> # Create new environment for time series
> ratesenv <- new.env()
> # Download time series for symbolv into ratesenv
> quantmod::getSymbols(symbolv, env=ratesenv, src="FRED")
> # Remove NA values in ratesenv
> sapply(ratesenv, function(x) sum(is.na(x)))
> sapply(ls(ratesenv), function(namev) {
+   assign(x=namev, value=na.omit(get(namev, ratesenv)),
+     env=ratesenv)
+   return(NULL)
+ }) ## end sapply
> sapply(ratesenv, function(x) sum(is.na(x)))
> # Get class of all objects in ratesenv
> sapply(ratesenv, class)
> # Get class of all objects in R workspace
> sapply(ls(), function(namev) class(get(namev)))
> # Save the time series environment into a binary .RData file
> save(ratesenv, file="/Users/jerzy/Develop/lecture_slides/data/rates_data.RData")
```



```
> # Get class of time series object DGS10
> class(get(x="DGS10", env=ratesenv))
> # Another way
> class(ratesenv$DGS10)
> # Get first 6 rows of time series
> head(ratesenv$DGS10)
> # Plot dygraphs of 10-year Treasury rate
> dygraphs::dygraph(ratesenv$DGS10, main="10-year Treasury Rate") %>%
+   dyOptions(colors="blue", strokeWidth=2)
> # Plot 10-year constant maturity Treasury rate
> x11(width=6, height=5)
> par(mar=c(2, 2, 0, 0), oma=c(0, 0, 0, 0))
> chart_Series(ratesenv$DGS10["1990/",], name="10-year Treasury Rate")
```


Treasury Yield Curve

The *yield curve* is a vector of interest rates at different maturities, on a given date.

The *yield curve* shape changes depending on the economic conditions: in recessions rates drop and the curve flattens, while in expansions rates rise and the curve steepens.

```
> # Load constant maturity Treasury rates
> load(file="/Users/jerzy/Develop/lecture_slides/data/rates_data.RData")
> # Get most recent yield curve
> ycnow <- eapply(ratesenv, xts::last)
> class(ycnow)
> ycnow <- do.call(cbind, ycnow)
> # Check if 2020-03-25 is not a holiday
> date2020 <- as.Date("2020-03-25")
> weekdays(date2020)
> # Get yield curve from 2020-03-25
> yc2020 <- eapply(ratesenv, function(x) x[date2020])
> yc2020 <- do.call(cbind, yc2020)
> # Combine the yield curves
> ycurves <- c(yc2020, ycnow)
> # Rename columns and rows, sort columns, and transpose into matrix
> colnames(ycurves) <- substr(colnames(ycurves), start=4, stop=11)
> ycurves <- ycurves[, order(as.numeric(colnames(ycurves)))]
> colnames(ycurves) <- paste0(colnames(ycurves), "yr")
> ycurves <- t(ycurves)
> colnames(ycurves) <- substr(colnames(ycurves), start=1, stop=4)
```

Yield Curves in 2020 and 2025




```
> # Plot using matplot()
> colorv <- c("blue", "red")
> matplot(ycurves, main="Yield Curves in 2020 and 2023", xaxt="n",
+ type="l", xlab="maturity", ylab="yield", col=colorv)
> # Add x-axis
> axis(1, seq_along(rownames(ycurves)), rownames(ycurves))
> # Add legend
> legend("topleft", legend=colnames(ycurves), y.intersp=0.5,
+ bty="n", col=colorv, lty=1, lwd=6, inset=0.05, cex=1.0)
```

Treasury Yield Curve Over Time

The *yield curve* changes shape dramatically depending on the economic conditions: in recessions rates drop and the curve flattens, while in expansions rates rise and the curve steepens.

```
> # Load constant maturity Treasury rates
> load(file="/Users/jerzy/Develop/lecture_slides/data/rates_data.RD")
> # Get end-of-year dates since 2006
> datev <- xts::endpoints(ratesenv$DGS1["2006/"], on="years")
> datev <- zoo::index(ratesenv$DGS1["2006/"][datev])
> # Create time series of end-of-year rates
> ycurves <- eapply(ratesenv, function(ratev) ratev[datev])
> ycurves <- rutils::do_call(cbind, ycurves)
> # Rename columns and rows, sort columns, and transpose into matrix
> colnames(ycurves) <- substr(colnames(ycurves), start=4, stop=11)
> ycurves <- ycurves[, order(as.numeric(colnames(ycurves))))]
> colnames(ycurves) <- paste0(colnames(ycurves), "yr")
> ycurves <- t(ycurves)
> colnames(ycurves) <- substr(colnames(ycurves), start=1, stop=4)
> # Plot matrix using plot.zoo()
> colorv <- colorRampPalette(c("red", "blue"))(NCOL(ycurves))
> plot.zoo(ycurves, main="Yield Curve Since 2006", lwd=3, xaxt="n"
+   plot.type="single", xlab="maturity", ylab="yield", col=colorv)
> # Add x-axis
> axis(1, seq_along(rownames(ycurves)), rownames(ycurves))
> # Add legend
> legend("topleft", legend=colnames(ycurves), y.intersp=0.5,
+   bty="n", col=colorv, lty=1, lwd=4, inset=0.05, cex=0.8)
```

Yield curve since 2006



```
> # Alternative plot using matplot()
> matplot(ycurves, main="Yield curve since 2006", xaxt="n", lwd=3,
+   type="l", xlab="maturity", ylab="yield", col=colorv)
> # Add x-axis
> axis(1, seq_along(rownames(ycurves)), rownames(ycurves))
> # Add legend
> legend("topleft", legend=colnames(ycurves), y.intersp=0.5,
+   bty="n", col=colorv, lty=1, lwd=4, inset=0.05, cex=0.8)
```

Nelson-Siegel Yield Curve Model

The Nelson-Siegel model of the *yield curve* is given by:

$$y(\tau) = \beta_0 + \beta_1 \left(\frac{1 - e^{-\lambda\tau}}{\lambda\tau} \right) + \beta_2 \left(\frac{1 - e^{-\lambda\tau}}{\lambda\tau} - e^{-\lambda\tau} \right)$$


Where $y(\tau)$ is the yield at the maturity τ and λ is the time decay factor.

The first term β_0 represents the long-term level of interest rates, the second term β_1 represents the slope, and the third term β_2 represents the curvature of the yield curve.

For very short maturities ($\tau \rightarrow 0$), the yield is equal to $y(0) = \beta_0 + \beta_1 + \beta_2$.

For very long maturities ($\tau \rightarrow \infty$), the yield is equal to $y(\infty) = \beta_0$

Nelson-Siegel Yield Curve



```
> # Define the Nelson-Siegel model
> slopet <- function(tau, lambda) (1 - exp(-lambda * tau)) / (lambda)
> curvet <- function(tau, lambda) (1 - exp(-lambda * tau)) / (lambda)
> yns <- function(tau, beta0, beta1, beta2, lambda) {
+   return(beta0 + beta1 * slopet(tau, lambda) + beta2 * curvet(tau, lambda))
+ } ## end yns
> # Plot the Nelson-Siegel components
> tau <- seq(0.1, 30, by=0.2)
> lambda <- 0.2
> beta0 <- 1.0; beta1 <- -1.5; beta2 <- 3.0
> matplot(tau, cbind(beta0, beta1 * slopet(tau, lambda), beta2 * curvet(tau, lambda)),
+   type="l", lwd=3, lty=1,
+   col=c("green", "blue", "red"),
+   main="Nelson-Siegel Yield Curve Components",
+   xlab="Maturity (years)", ylab="Yield (%)")
> legend("bottomright", bty="n", col=c("green", "blue", "red"),
+   legend=c("beta0", "beta1", "beta2"))

> # Plot the Nelson-Siegel yield curve
> yieldv <- yns(tau, beta0, beta1, beta2, lambda)
> plot(tau, yieldv, type="l", lwd=3, col="blue",
+   main="Nelson-Siegel Yield Curve",
+   xlab="Maturity (years)", ylab="Yield (%)")
> # Plot the Nelson-Siegel yield curve for different values of lambda
> lambdav <- c(0.1, 0.5, 1.0, 2.0)
> yieldm <- sapply(lambdav, function(lambda) {
+   yns(tau, beta0, beta1, beta2, lambda)
+ }) ## end lapply
> matplot(tau, yieldm, type="l", lwd=3, lty=1,
+   col=rainbow(length(lambdav)),
+   main="Nelson-Siegel Yield Curve for Different Lambda",
+   xlab="Maturity (years)", ylab="Yield (%)")
> legend("bottomright", legend=paste0("lambda=", lambdav),
```

Package *DEoptim* for Global Optimization

The function `DEoptim()` from package *DEoptim* performs *global* optimization using the *Differential Evolution* algorithm.

Differential Evolution is a genetic algorithm which evolves a population of solutions over several generations:

<https://link.springer.com/content/pdf/10.1023/A:1008202821328.pdf>

The first generation of solutions is selected randomly.

Each new generation is obtained by combining the best solutions from the previous generation.

The *Differential Evolution* algorithm is well suited for very large multi-dimensional optimization problems, such as portfolio optimization.

Gradient optimization methods are more efficient than *Differential Evolution* for smooth objective functions with no local minima.

A limitation of *Differential Evolution* is that it doesn't estimate the standard errors of the parameters.

```
> # Rastrigin function with vector argument for optimization
> rastrigin <- function(vecv, param=25) {
+   sum(vecv^2 - param*cos(vecv))
+ } ## end rastrigin
> vecv <- c(pi/6, pi/6)
> rastrigin(vecv=vecv)
> library(DEoptim)
> # Optimize rastrigin using DEoptim
> optim1 <- DEoptim(rastrigin,
+   upper=c(6, 6), lower=c(-6, -6),
+   DEoptim.control(trace=FALSE, itermx=50))
> # Optimal parameters and value
> optim1$optim$bestmem
> rastrigin(optim1$optim$bestmem)
> summary(optim1)
> plot(optim1)
```


Calibrating the Nelson-Siegel Model

The Nelson-Siegel model can be calibrated to the actual yields using *global* optimization, to avoid local minima.

The function `DEoptim()` from package *DEoptim* performs *global* optimization using the *Differential Evolution* algorithm.

```
> # Extract numeric maturities
> yieldv <- ycurves[, "2025"]
> tau <- as.numeric(sub("yr", "", names(yieldv)))
> # Define the Nelson-Siegel model objective function
> objfun <- function(params, tau, yieldv) {
+   beta0 <- params[1]; beta1 <- params[2]; beta2 <- params[3]; lam1
+   yieldf <- yns(tau, beta0, beta1, beta2, lambda)
+   return(sum((yieldv - yieldf)^2))
+ } ## end objfun
> # Run DEoptim
> set.seed(1234)
> library(DEoptim)
> optim1 <- DEoptim(fn = objfun,
+   tau=tau, yieldv=yieldv,
+   lower = c(-20, -20, -20, -20),
+   upper = c(20, 20, 20, 20),
+   control = DEoptim.control(trace=FALSE, storepopfrom = 1, iterm:
> # The optimal parameters
> optimp <- optim1$optim$bestmem
> beta0 <- optimp[1]; beta1 <- optimp[2]; beta2 <- optimp[3]; lamb
```

Nelson-Siegel and Actual Yield Curve



```
> # Plot the Nelson-Siegel yield curve
> yieldv <- yns(tau, beta0, beta1, beta2, lambda)
> matplot(tau, cbind(yieldv, ycurves[, "2025"]),
+   type="l", lty=1, lwd=3, col=c("red", "blue"),
+   main="Nelson-Siegel and Actual Yield Curve",
+   xlab="Maturity (years)", ylab="Yield (%)")
> legend("bottomright", legend=c("Nelson-Siegel", "Actual"),
+   bty="n", col=c("red", "blue"),
+   lty=1, lwd=3, inset=0.05, cex=1.0)
```


Covariance Matrix of Interest Rates

The covariance matrix $\mathbb{C}$, of the interest rate matrix $\mathbf{r}$ is given by:

$$\mathbb{C} = \frac{(\mathbf{r} - \bar{\mathbf{r}})^T (\mathbf{r} - \bar{\mathbf{r}})}{n - 1}$$

```
> # Extract rates from ratesenv
> symbolv <- c("DGS1", "DGS2", "DGS5", "DGS10", "DGS20")
> ratem <- mget(symbolv, envir=ratesenv)
> ratem <- rutils::do_call(cbind, ratem)
> ratem <- zoo::na.locf(ratem, na.rm=FALSE)
> ratem <- zoo::na.locf(ratem, fromLast=TRUE)
> # Calculate daily percentage rates changes
> retp <- rutils::diffit(log(ratem))
> # Center (de-mean) the returns
> retp <- lapply(retp, function(x) {x - mean(x)})
> retp <- rutils::do_call(cbind, retp)
> sapply(retp, mean)
> # Covariance and Correlation matrices of Treasury rates
> covmat <- cov(retp)
> cormat <- cor(retp)
> # Reorder correlation matrix based on clusters
> library(corrplot)
> ordern <- corrMatOrder(cormat, order="hclust",
+   hclust.method="complete")
> cormat <- cormat[ordern, ordern]
```

Correlation of Treasury Rates



```
> # Plot the correlation matrix
> x11(width=6, height=6)
> colorv <- colorRampPalette(c("red", "white", "blue"))
> corrplot(cormat, title=NA, tl.col="black",
+   method="square", col=colorv(NCOL(cormat)), tl.cex=0.8,
+   cl.offset=0.75, cl.cex=0.7, cl.align.text="l", cl.ratio=0.25)
> title("Correlation of Treasury Rates", line=1)
> # Draw rectangles on the correlation matrix plot
> corrRect.hclust(cormat, k=NROW(cormat) %/% 2,
+   method="complete", col="red")
```


Principal Component Vectors

Principal components are linear combinations of the k return vectors $\mathbf{r}_i$:

$$\mathbf{pc}_j = \sum_{i=1}^k w_{ij} \mathbf{r}_i$$


Where $\mathbf{w}_j$ is a vector of weights (loadings) of the *principal component* j , with $\mathbf{w}_j^T \mathbf{w}_j = 1$.

The weights $\mathbf{w}_j$ are chosen to maximize the variance of the *principal components*, under the condition that they are orthogonal to:

$$\mathbf{w}_j = \arg \max \left\{ \mathbf{pc}_j^T \mathbf{pc}_j \right\}$$

$$\mathbf{pc}_i^T \mathbf{pc}_j = 0 \quad (i \neq j)$$

```
> # Create initial vector of portfolio weights
> nweights <- NROW(symbolv)
> weightv <- rep(1/sqrt(nweights), nweights)
> names(weightv) <- symbolv
> # Objective function equal to minus portfolio variance
> objfun <- function(weightv, retp) {
+   retp <- retp %*% weightv
+   -1e7*var(retp) + 1e7*(1 - sum(weightv*weightv))^2
+ } ## end objfun
> # Objective function for equal weight portfolio
> objfun(weightv, retp)
> # Compare speed of vector multiplication methods
> library(microbenchmark)
> summary(microbenchmark(
+   transp=t(retp) %*% retp,
+   sumv=sum(retp*retp),
```



```
> # Find weights with maximum variance
> optim1 <- optim(par=weightv,
+   fn=objfun,
+   retp=retp,
+   method="L-BFGS-B",
+   upper=rep(5.0, nweights),
+   lower=rep(-5.0, nweights))
> # Optimal weights and maximum variance
> weights1 <- optim1$par
> objfun(weights1, retp)
> # Plot first principal component loadings
> x11(width=6, height=5)
> par(mar=c(3, 3, 2, 1), oma=c(0, 0, 0, 0), mgp=c(2, 1, 0))
> barplot(weights1, names.arg=names(weights1),
+   xlab="", ylab="", main="First Principal Component Loadings")
```


Higher Order Principal Components

The *second principal component* can be calculated by maximizing its variance, under the constraint that it must be orthogonal to the *first principal component*.

Similarly, higher order *principal components* can be calculated by maximizing their variances, under the constraint that they must be orthogonal to all the previous *principal components*.

The number of principal components is equal to the dimension of the covariance matrix.

```
> # pc1 weights and returns
> pc1 <- drop(retp %*% weights1)
> # Redefine objective function
> objfun <- function(weightv, retp) {
+   retp <- retp %*% weightv
+   -1e7*var(retp) + 1e7*(1 - sum(weightv^2))^2 +
+   1e7*sum(weights1*weightv)^2
+ } ## end objfun
> # Find second principal component weights
> optiml <- optim(par=weightv,
+   fn=objfun,
+   retp=retp,
+   method="L-BFGS-B",
+   upper=rep(5.0, nweights),
+   lower=rep(-5.0, nweights))
```



```
> # pc2 weights and returns
> weights2 <- optiml$par
> pc2 <- drop(retp %*% weights2)
> sum(pc1*pc2)
> # Plot second principal component loadings
> barplot(weights2, names.arg=names(weights2),
+   xlab="", ylab="", main="Second Principal Component Loadings")
```

Eigenvalues of the Covariance Matrix

The portfolio variance: $\mathbf{w}^T \mathbb{C} \mathbf{w}$ can be maximized under the *quadratic* weights constraint $\mathbf{w}^T \mathbf{w} = 1$, by maximizing the *Lagrangian* $\mathcal{L}$:

$$\mathcal{L} = \mathbf{w}^T \mathbb{C} \mathbf{w} - \lambda (\mathbf{w}^T \mathbf{w} - 1)$$

Where λ is a *Lagrange multiplier*.


The maximum variance portfolio weights can be found by differentiating $\mathcal{L}$ with respect to $\mathbf{w}$ and setting it to zero:

$$\mathbb{C} \mathbf{w} = \lambda \mathbf{w}$$

The above is the *eigenvalue* equation of the covariance matrix $\mathbb{C}$, with the optimal weights $\mathbf{w}$ forming an *eigenvector*, and λ is the *eigenvalue* corresponding to the *eigenvector* $\mathbf{w}$.

The *eigenvalues* are the variances of the *eigenvectors*, and their sum is equal to the sum of the return variances:

$$\sum_{i=1}^k \lambda_i = \frac{1}{1-k} \sum_{i=1}^k \mathbf{r}_i^T \mathbf{r}_i$$



```
> eigend <- eigen(covmat)
> eigend$eigenvectors
> # Compare with optimization
> all.equal(sum(diag(covmat)), sum(eigend$values))
> all.equal(abs(eigend$eigenvectors[, 1]), abs(weights1), check.attributes=FALSE)
> all.equal(abs(eigend$eigenvectors[, 2]), abs(weights2), check.attributes=FALSE)
> all.equal(eigend$values[1], var(pc1), check.attributes=FALSE)
> all.equal(eigend$values[2], var(pc2), check.attributes=FALSE)
> # Eigenvalue equations are satisfied approximately
> (covmat %*% weights1) / weights1 / var(pc1)
> (covmat %*% weights2) / weights2 / var(pc2)
> # Plot eigenvalues
> barplot(eigend$values, names.arg=paste0("PC", 1:nweights),
+   las=3, xlab="", ylab="", main="Principal Component Variances")
```

Principal Component Analysis Versus Eigen Decomposition

Principal Component Analysis (PCA) is equivalent to the *eigen decomposition* of either the correlation or the covariance matrix.

If the input time series *are* scaled, then *PCA* is equivalent to the eigen decomposition of the *correlation matrix*.

If the input time series *are not* scaled, then *PCA* is equivalent to the eigen decomposition of the *covariance matrix*.

Scaling the input time series improves the accuracy of the *PCA dimension reduction*, allowing a smaller number of *principal components* to more accurately capture the data contained in the input time series.

The function `prcomp()` performs *Principal Component Analysis* on a matrix of data (with the time series as columns), and returns the results as a list of class `prcomp`.

The `prcomp()` argument `scale=TRUE` specifies that the input time series should be scaled by their standard deviations.

```
> # Eigen decomposition of correlation matrix
> eigend <- eigen(cormat)
> # Perform PCA with scaling
> pcad <- prcomp(retp, scale=TRUE)
> # Compare outputs
> all.equal(eigend$values, pcad$sdev^2)
> all.equal(abs(eigend$vectors), abs(pcad$rotation),
+   check.attributes=FALSE)
> # Eigen decomposition of covariance matrix
> eigend <- eigen(covmat)
> # Perform PCA without scaling
> pcad <- prcomp(retp, scale=FALSE)
> # Compare outputs
> all.equal(eigend$values, pcad$sdev^2)
> all.equal(abs(eigend$vectors), abs(pcad$rotation),
+   check.attributes=FALSE)
```

Principal Component Analysis of the Yield Curve

Principal Component Analysis (PCA) can be used for *dimension reduction*, to explain a large number of correlated time series as linear combinations of a smaller number of principal component time series.


The input time series are often scaled by their standard deviations, to improve the accuracy of *PCA dimension reduction*, so that more information is retained by the first few *principal component* time series.

If the input time series are not scaled, then *PCA* analysis is equivalent to the *eigen decomposition* of the covariance matrix, and if they are scaled, then *PCA* analysis is equivalent to the *eigen decomposition* of the correlation matrix.

The function `prcomp()` performs *Principal Component Analysis* on a matrix of data (with the time series as columns), and returns the results as a list of class `prcomp`.

The `prcomp()` argument `scale=TRUE` specifies that the input time series should be scaled by their standard deviations.

Scree Plot: Volatilities of Principal Components of Treasury rates



A *scree plot* is a bar plot of the volatilities of the *principal components*.

```
> # Perform principal component analysis PCA
> pcam <- prcomp(retp, scale=TRUE)
> # Plot standard deviations
> barplot(pcam$sdev, names.arg=colnames(pcam$rotation),
+   las=3, xlab="", ylab="",
+   main="Scree Plot: Volatilities of Principal Components
+   of Treasury rates")
```

Yield Curve Principal Component Loadings (Weights)

Principal component loadings are the weights of portfolios which have mutually orthogonal returns.


The *principal component* portfolios represent the different orthogonal modes of the data variance.

The first *principal component* of the *yield curve* is the correlated movement of all rates up and down.

The second *principal component* is *yield curve* steepening and flattening.

The third *principal component* is the *yield curve* butterfly movement.

```
> # Calculate principal component loadings (weights)
> pcad$rotation
> # Plot loading barplots in multiple panels
> par(mfrow=c(3,2))
> par(mar=c(3.5, 2, 2, 1), oma=c(0, 0, 0, 0))
> for (ordern in 1:NCOL(pcad$rotation)) {
+   barplot(pcad$rotation[, ordern], las=3, xlab="", ylab="", main=
+ title(paste0("PC", ordern), line=-2.0, col.main="red")
+ } ## end for
```




Yield Curve Principal Component Time Series

The time series of the *principal components* can be calculated by multiplying the loadings (weights) times the original data.

The *principal component* time series have mutually orthogonal returns.

Higher order *principal components* are gradually less volatile.

```
> # Standardize (center and scale) the returns
> retp <- lapply(retp, function(x) {(x - mean(x))/sd(x)})
> retp <- rutils::do_call(cbind, retp)
> sapply(retp, mean)
> sapply(retp, sd)
> # Calculate principal component time series
> retpcac <- retp %>% pcad$rotation
> all.equal(pcad$x, retpcac, check.attributes=FALSE)
> # Calculate products of principal component time series
> round(t(retpcac) %>% retpcac, 2)
> # Coerce to xts time series
> retpcac <- xts(retpcac, order.by=zoo::index(retp))
> retpcac <- cumsum(retpcac)
> # Plot principal component time series in multiple panels
> par(mfrow=c(3,2))
> par(mar=c(2, 2, 0, 1), oma=c(0, 0, 0, 0))
> rangev <- range(retpcac)
> for (ordern in 1:NCOL(retpcac)) {
+   plot.zoo(retpcac[, ordern], ylim=rangev, xlab="", ylab="")
+   title(paste0("PC", ordern), line=-1, col.main="red")
+ } ## end for
```



Inverting Principal Component Analysis

The original time series can be calculated *exactly* from the time series of all the *principal components*, by inverting the loadings matrix.

The function `solve()` solves systems of linear equations, and also inverts square matrices.

```
> # Invert all the principal component time series
> retpca <- retp %*% pcad$rotation
> solved <- retpca %*% solve(pcad$rotation)
> all.equal(coredata(retp), solved)
```



Dimension Reduction Using Principal Component Analysis

The original time series can be calculated *approximately* from just the first few *principal components*, which demonstrates how *PCA* can be used for *dimension reduction*.

A popular rule of thumb is to use the *principal components* with the largest variances, which sum up to 80% of the total variance of returns.

The *Kaiser-Guttman* rule uses only *principal components* with variance greater than 1.

```
> # Invert first 3 principal component time series
> solved <- retpcal[, 1:3] %*% solve(pcad$rotation)[1:3, ]
> solved <- xts::xts(solved, zoo::index(retp))
> solved <- cumsum(solved)
> retc <- cumsum(retp)
> # Plot the solved returns
> par(mfrow=c(3,2))
> par(mar=c(2, 2, 0, 1), oma=c(0, 0, 0, 0))
> for (symbol in symbolv) {
+   plot.zoo(cbind(retc[, symbol], solved[, symbol]),
+   plot.type="single", col=c("black", "blue"), xlab="", ylab="")
+   legend(x="topleft", bty="n", y.intersp=0.5,
+   legend=paste0(symboln, c("", " solved")),
+   title=NULL, inset=0.0, cex=1.0, lwd=6,
+   lty=1, col=c("black", "blue"))
+ } ## end for
```



Conservation of Variance in PCA

The matrix of *principal component* returns $\mathbf{pc}$ is equal to the matrix $\mathbf{w}$ of the PCA weights (loadings) times the matrix of the return vectors $\mathbf{r}$:

$$\mathbf{pc} = \mathbf{w} \mathbf{r}$$

Where the PCA weights $\mathbf{w}^T \mathbf{w} = \mathbb{1}$ are orthonormal.

So the sum of the PCA variances is equal to:

$$\mathbf{pc}^T \mathbf{pc} = (\mathbf{w} \mathbf{r})^T \mathbf{w} \mathbf{r} = \mathbf{r}^T (\mathbf{w}^T \mathbf{w}) \mathbf{r} = \mathbf{r}^T \mathbf{r}$$

So the total variance is conserved in PCA: the sum of the PCA variances is equal to the sum of the factor variances.

```
> # Calculate the PC time series
> pcad <- prcomp(retp, scale=FALSE)
> retpca <- retp %*% pcad$rotation
> # Compare the total variances
> all.equal(sum(apply(retpca, 2, var)), sum(apply(retp, 2, var)))
```

Calibrating Yield Curve Using Package *RQuantLib*

The package *RQuantLib* is an interface to the *QuantLib* open source C/C++ library for quantitative finance, mostly designed for pricing fixed-income instruments and options.

The function `DiscountCurve()` calibrates a *zero coupon yield curve* from *money market rates*, *Eurodollar futures*, and *swap rates*.

The function `DiscountCurve()` interpolates the *zero coupon rates* into a vector of dates specified by the `times` argument.

```
> library(RQuantLib) ## Load RQuantLib
> # Specify curve parameters
> curvep <- list(tradeDate=as.Date("2018-01-17"),
+               settleDate=as.Date("2018-01-19"),
+               dt=0.25,
+               interpWhat="discount",
+               interpHow="loglinear")
> # Specify market data: prices of FI instruments
> pricev <- list(d3m=0.0363,
+               fut1=96.2875,
+               fut2=96.7875,
+               fut3=96.9875,
+               fut4=96.6875,
+               s5y=0.0443,
+               s10y=0.05165,
+               s15y=0.055175)
> # Specify dates for calculating the zero rates
> datev <- seq(0, 10, 0.25)
> # Specify the evaluation (as of) date
> setEvaluationDate(as.Date("2018-01-17"))
> # Calculate the zero rates
> ratev <- DiscountCurve(params=curvep, tsQuotes=pricev, times=datev)
> # Plot the zero rates
> x11()
> plot(x=ratev$zerorates, t="1", main="zerorates")
```

Financial and Commodity Futures Contracts

The underlying assets delivered in *commodity futures* contracts are commodities, such as grains (corn, wheat), or raw materials and metals (oil, aluminum).

The underlying assets delivered in *financial futures* contracts are financial assets, such as stocks, bonds, and currencies.

Many futures contracts use cash settlement instead of physical delivery of the asset.

Futures contracts on different underlying assets can have quarterly, monthly, or even weekly expiration dates.

The front month futures contract is the contract with the closest expiration date to the current date.

Symbols of futures contracts are obtained by combining the contract code with the month code and the year.

For example, *ESM9* is the symbol for the *S&P500* index E-mini futures expiring in June 2019.

Futures contract	Code
S&P500 index	ES
10yr Treasury	ZN
VIX index	VX
Gold	GC
Oil	CL
Euro FX	EC
Swiss franc	SF
Japanese Yen	JY

Month	Code
January	F
February	G
September	H
April	J
May	K
June	M
July	N
August	Q
September	U
October	V
November	X
December	Z

Interactive Brokers provides more information about futures contracts:

[IB Contract and Symbol Database](#)

[IB Traded Products](#)

List of [Popular Futures Contracts](#).

E-mini Futures Contracts

E-mini futures are contracts with smaller notionals and tick values, which are more suitable for retail investors.

For example, the *QM E-mini oil future* notional is 500 barrels, while the standard *CL oil future* notional is 1,000 barrels.

The tick value is the change in the dollar value of the futures contract due to a one tick change in the underlying price.

For example, the tick value of the *ES E-mini S&P500* future is \$12.50, and one tick is 0.25.

So if the *S&P500* index changes by one tick (0.25), then the value of a single *ES* E-mini contract changes by \$12.50, while the standard *SP* contract value changes by \$62.5.

The *ES E-mini S&P500 futures* trade almost continuously 24 hours per day, from 6:00 PM Eastern Time (ET) on Sunday night to 5:00 PM Friday night (with a trading halt between 4:15 and 4:30 PM ET each day).

Futures contract	Standard	E-mini
S&P500 index	SP	ES
10yr Treasury	ZN	ZN
VIX index	VX	delisted
Gold	GC	YG
Oil	CL	QM
Euro FX	EC	E7
Swiss franc	SF	MSF
Japanese Yen	JY	J7

S&P500 Futures Prices


The S&P500 E-mini futures prices in February and March 2025 followed a typical market pattern: the trading volumes rose as the prices were dropping.

The function `data.table::fread()` reads .csv files over five times faster than the function `read.csv()`!

The function `as.Date()` parses character strings and coerces numeric objects into Date objects.

The function `as.POSIXct.numeric()` coerces a numeric value representing the *moment of time* into a POSIXct *date-time*, equal to the *clock time* in the local *time zone*.

```
> # Load ESH5 March 2025 E-mini futures prices
> fileN <- "/Users/jerzy/Develop/data/futures/ESH5.csv"
> ESH5 <- data.table::fread(fileN)
> # Coerce ESH5 into xts series
> ESH5$Date <- as.Date(ESH5$Date, format="%m/%d/%y")
> ESH5 <- data.table::as.xts.data.table(ESH5)
> colnames(ESH5) <- c("Open", "High", "Low", "Close", "Volume")
> tail(ESH5)
> # Candlestick plot of the ESH5 prices
> themev <- chart_theme()
> themev$col$up.col <- "green"
> quantmod::chart_Series(x=ESH5["2025-01/"], TA="add_Vo()",
+   theme=themev, name="March 2025 E-mini Futures Prices")
```



```
> # Plot dygraph candlestick of ESH5 prices
> dygraphs::dygraph(ESH5["2025-01/", 1:4],
+   main="March 2025 E-mini Futures Prices") %>%
+   dyCandlestick() %>% dyLegend(show="onmouseover")
> # Plot dygraph barplot of ESH5 volumes
> dygraphs::dygraph(ESH5["2025-01/", 5],
+   main="March 2025 E-mini Futures Volumes") %>%
+   dyBarChart() %>% dyOptions(colors="blue", strokeWidth=3)
```


Consecutive Contract Futures Volumes

The trading volumes of a futures contract quickly drop shortly before its expiration, and at the same time, the volumes of the successive contract increase.

The contract with the highest trading volume is usually considered the most liquid contract.

ESZ4 is the December 2024 E-mini futures contract, and *ESH5* is the March 2025 E-mini futures contract.

```
> # Load ESZ4 December 2024 E-mini futures prices
> filen <- "/Users/jerzy/Develop/data/futures/ESZ4.csv"
> ESZ4 <- data.table::fread(filen)
> ESZ4$Date <- as.Date(ESZ4$Date, format="%m/%d/%y")
> ESZ4 <- data.table::as.xts.data.table(ESZ4)
> colnames(ESZ4) <- c("Open", "High", "Low", "Close", "Volume")
```



```
> # Plot last month of ESZ4 and ESH5 volume data
> endd <- end(ESZ4)
> startd <- (endd - 30)
> volumm <- cbind(Vo(ESZ4), Vo(ESH5))[paste0(startd, "/", endd)]
> colnames(volumm) <- c("ESZ4", "ESH5")
> dygraphs::dygraph(volumm,
+   main="Volumes of ESZ4 and ESH5 futures") %>%
+   dyOptions(colors=c("blue", "green"), strokeWidth=3) %>%
+   dyLegend(show="always", width=300)
```

Chaining Together Futures Prices

Chaining futures means splicing together prices from several consecutive futures contracts.

A continuous futures contract is a time series of prices obtained by chaining together prices from consecutive futures contracts.

When the next contract becomes more liquid, then the continuous futures price is rolled over to that contract.

The price of the continuous futures is equal to the most liquid contract times a scaling factor.

Futures contracts with different maturities (expiration dates) trade at different prices because of the futures curve, which causes price jumps between consecutive futures contracts.

The old contract price is multiplied by a scaling factor after that contract is rolled, to remove price jumps.

So the continuous futures prices are not equal to the past futures prices.

```
> # Find date when ESH5 volume exceeds ESZ4
> xeed <- (volumm[, "ESH5"] > volumm[, "ESZ4"])
> indeks <- match(TRUE, xeed)
> # indeks <- min(which(xeed))
> indeks <- zoo::index(volumm[indeks])
> # Scale the ESZ4 prices
> ratiop <- as.numeric(C1(ESH5[indeks]))/as.numeric(C1(ESZ4[indeks]))
> ESZ4[, 1:4] <- ratiop*ESZ4[, 1:4]
```



```
> # Calculate the continuous futures prices
> pricev <- rbind(ESZ4[paste0("/", indeks-1)],
+               ESH5[paste0(indeks, "/")])
> # Candlestick plot of the continuous E-mini futures prices
> themev <- chart_theme()
> themev$col$up.col <- "green"
> quantmod::chart_Series(x=pricev["2024-09/", TA="add_Vo()"],
+   theme=themev, name="Continuous E-mini Futures Prices")
> # dygraph the continuous E-mini futures prices
> dygraphs::dygraph(quantmod::C1(pricev["2024-09/"]),
+   main="Continuous E-mini Futures Prices") %>%
+   dyOptions(colors="blue", strokeWidth=2)
```


VIX Volatility Index

The *VIX* Volatility Index is an average of the implied volatilities of options on the *S&P500* Index (SPX).

The *VIX* index is an estimate of the *future* stock market volatility that is expected (anticipated) by investors.

The *VIX* index is not a directly tradable asset, but it can be traded using *VIX* futures.

The *CBOE* (Chicago Board Options Exchange) calculates the *VIX* index and provides free daily historical data for the *VIX* and other volatility indices.

The *CBOE* also provides free daily historical *VIX* futures prices.

The *FRED* database also provides historical data for the *VIX* index and for other volatility indices.

VIX Index



```
> # Download VIX index data from CBOE
> vixindex <- data.table::fread("https://cdn.cboe.com/api/global/u
> class(vixindex)
> dim(vixindex)
> tail(vixindex)
> sapply(vixindex, class)
> vixindex <- xts(vixindex[, -1],
+   order.by=as.Date(vixindex$V1, format="%m/%d/%Y"))
> colnames(vixindex) <- c("Open", "High", "Low", "Close")
> # Save the VIX data to binary file
> load(file="/Users/jerzy/Develop/data/vix/vix_data.RData")
> # vixenv <- new.env()
> ls(vixenv)
> vixenv$vixindex <- vixindex
> save(vixenv, file="/Users/jerzy/Develop/data/vix/vix_data.RData")
```

```
> # Plot dygraph
> dygraphs::dygraph(vixindex["2019/"], main="VIX Index") %>%
+   dyCandlestick()
> # Plot VIX OHLC data in x11 window
> chart_Series(x=vixindex["2018"], name="VIX Index")
```

VIX Futures Contracts

VIX futures are cash-settled futures contracts on the VIX Index.

The most liquid VIX futures are with monthly expiration dates ([CBOE Expiration Calendar](#)), but weekly VIX futures are also traded.

These are the [VIX Futures Monthly Expiration Dates](#) from 2004 to 2019.

VIX futures are traded on the CFE (CBOE Futures Exchange):

<http://cfe.cboe.com/>

<http://www.cboe.com/vix>

VIX Contract Specifications:

[VIX Contract Specifications](#)

[VIX Expiration Calendar](#)

Standard and Poor's explains the methodology of the [VIX Futures Indices](#)

```
> # Read CBOE monthly futures expiration dates
> datev <- read.csv(file="/Users/jerzy/Develop/data/vix_data/vix_data.csv")
> datev <- as.Date(datev[, 1])
> yearv <- format(datev, format="%Y")
> yearv <- substring(yearv, 4)
> # Monthly futures contract codes
> codes <-
+   c("F", "G", "H", "J", "K", "M",
+     "N", "Q", "U", "V", "X", "Z")
> symbolv <- paste0("VX", codes, yearv)
> datev <- as.data.frame(datev)
> colnames(datev) <- "exp_dates"
> rownames(datev) <- symbolv
> # Write dates to CSV file, with row names
> write.csv(datev, row.names=TRUE,
+   file="/Users/jerzy/Develop/data/vix_data/vix_futures.csv")
> # Read back CBOE futures expiration dates
> datev <- read.csv(file="/Users/jerzy/Develop/data/vix_data/vix_futures.csv")
+   row.names=1)
> datev[, 1] <- as.Date(datev[, 1])
```

VIX Futures Curve

Futures contracts with different expiration dates trade at different prices, known as the *futures curve* (or *term structure*).

The *VIX* futures curve is similar to the interest rate *yield curve*, which displays yields at different bond maturities.

The *VIX* futures curve is not the same as the *VIX* index term structure.

More information about the *VIX* Index and the *VIX* futures curve:

- VIX Futures
- VIX Futures Data
- VIX Futures Curve
- VIX Index Term Structure

```
> # Load VIX futures data from binary file
> load(file="/Users/jerzy/Develop/data/vix_data/vix_data.RData")
> # Get all VIX futures for 2018 except January
> symbolv <- ls(vixenv)
> symbolv <- symbolv[grep(glob2rx("*8"), symbolv)]
> symbolv <- symbolv[2:9]
> # Specify dates for curves
> startd <- as.Date("2018-01-11")
> endd <- as.Date("2018-02-05")
> # Extract all VIX futures prices on the dates
> futcurves <- lapply(symbolv, function(symboln) {
+   xtsv <- get(x=symboln, envir=vixenv)
+   Cl(xtsv[c(startd, endd)])
+ }) ## end lapply
> futcurves <- rutils::do_call(cbind, futcurves)
> colnames(futcurves) <- symbolv
> futcurves <- t(coredata(futcurves))
> colnames(futcurves) <- c("Contango 01/11/2018",
+   "Backwardation 02/05/2018")
```


Contango and Backwardation of VIX Futures Curve

When prices are *low* then the futures curve is usually *upward sloping*, known as *contango*.

Futures prices are in *contango* most of the time.

When prices are *high* then the curve is usually *downward sloping*, known as *backwardation*.

```
> x11(width=7, height=5)
> par(mar=c(3, 2, 1, 1), oma=c(0, 0, 0, 0))
> plot(futcurves[, 1], type="l", lty=1, col="blue", lwd=3,
+      xaxt="n", xlab="", ylab="", ylim=range(futcurves),
+      main="VIX Futures Curves")
> axis(1, at=(1:NROW(futcurves)), labels=rownames(futcurves))
> lines(futcurves[, 2], lty=1, lwd=3, col="red")
> legend(x="topright", legend=colnames(futcurves),
+       inset=0.05, cex=1.0, bty="n", y.intersp=0.5,
+       col=c("blue", "red"), lwd=6, lty=1)
```



Futures Prices at Constant Maturity

A constant maturity futures price is the price of a hypothetical futures contract with an expiration date at a fixed number of days in the future.

Futures prices at a constant maturity can be calculated by interpolating the prices of contracts with neighboring expiration dates.

```
> # Load VIX futures data from binary file
> load(file="/Users/jerzy/Develop/data/vix_data/vix_data.RData")
> # Read CBOE futures expiration dates
> datev <- read.csv(file="/Users/jerzy/Develop/data/vix_data/vix_fut
+   row.names=1)
> symbolv <- rownames(datev)
> datev <- as.Date(datev[, 1])
> tday <- as.Date("2018-05-07")
> maturd <- (tday + 30)
> # Find neighboring futures contracts
> indeks <- match(TRUE, datev > maturd)
> frontd <- datev[indeks-1]
> backd <- datev[indeks]
> symbolf <- symbolv[indeks-1]
> symbolb <- symbolv[indeks]
> pricef <- get(x=symbolf, envir=vixenv)
> pricef <- as.numeric(Cl(pricef[tday]))
> priceb <- get(x=symbolb, envir=vixenv)
> priceb <- as.numeric(Cl(priceb[tday]))
> # Calculate the constant maturity 30-day futures price
> ratiop <- as.numeric(maturd - frontd)/as.numeric(backd - frontd)
> pricec <- (ratiop*priceb + (1-ratiop)*pricef)
```

VIX Futures Investing

The volatility index moves in the opposite direction to the underlying asset price.

An increase in the *VIX* index coincides with a drop in stock prices, and vice versa.

Taking a *long* position in *VIX* futures is similar to a *short* position in stocks, and vice versa.

There are several exchange-traded funds (*ETFs*) and exchange traded notes (*ETNs*) which are linked to *VIX* futures.

The *VXX* *ETN* provides the total return on the long *VIX* futures index.


VXX is short market risk because it's long *VIX* futures, and the *VIX* rises when stock prices drop.

The *SVXY* *ETF* provides the total return of short *VIX* futures contracts.

SVXY is long market risk because it's short *VIX* futures, and the *VIX* drops when stock prices rise.

Standard and Poor's explains the calculation of the **Total Return on VIX Futures Indices**.

```
> # Plot VIX and SVXY data in x11 window
> themev <- chart_theme()
> themev$col$line.col <- "blue"
> chart_Series(x=Cl(vixenv$vixindex["2007/"]),
+             theme=themev, name="VIX Index")
> chart_Series(x=Cl(rutils::etfenv$VTI["2007/"]),
+             theme=themev, name="VTI ETF")
```



VIX Crash on February 5th 2018

The SVXY and XIV ETFs rallied strongly after the financial crisis of 2008, so they became very popular with individual investors, and became very "crowded trades".


The SVXY and XIV ETFs had \$3.6 billion of assets at the beginning of 2018.

On February 5th 2018 the U.S. stock markets experienced a mini-crash, which was exacerbated by VIX futures short sellers.

As a result, the XIV ETF hit its termination event and its value dropped to zero:

Volatility Caused Stock Market Crash
XIV ETF Termination Event

```
> chart_Series(x=Cl(vixenv$vixindex["2017/2018"]),
+             theme=themev, name="VIX Index")
> chart_Series(x=Cl(rutils::etfenv$SVXY["2017/2018"]),
+             theme=themev, name="SVXY ETF")
```



draft: Types of Market Data

Fundamental company data summarizes the financial, economic, and regulatory state of a company.

Fundamental company data can be divided into several different types:

- Balance Sheet (assets and liabilities),
- Income Statement (profits and losses),
- Cash Flow Statement (current operating cash flows),
- Financial Ratios (performance and risk measures),

Financial ratios summarize the performance and risk measures of a company, and can be used for investment decisions, like value investing.

```
> # Read table of fundamental data into data frame
> dfame <- read.csv(file="/Users/jerzy/Develop/lecture_slides/data/
```

Balance.Sheet	Income.Statement	Cash.Flow.Statement	Financial.Ratios
Cash & Equivalents	Total Revenue	Cash from Operations	Earnings per Share
Investments	Cost of Goods Sold	Cash from Investments	Price/Earnings Ratio
Receivables	Sales, G and A	Cash from Financing	EBIT
Property & Equipment	Research & Development	FX Adjustment	EBITDA
Goodwill & Intangibles	Depreciation & Amortization	Capital Expenditure	Dividend Yield
Total Assets	Operating Expenses	Starting Cash	Gross Margin
Accounts Payable	Operating Income	Ending Cash	Profit Margin
Long-Term Debt	Investment Income		Free Cash Flow
Total Liabilities	Non-Operating Income		Asset Turnover
Common Stock	Pre-Tax Income		Inventory Turnover
			Receivable Turnover

Types of Fundamental Company Data

Fundamental company data summarizes the financial, economic, and regulatory state of a company.

Fundamental company data can be divided into several different types:

- Balance Sheet (assets and liabilities),
- Income Statement (profits and losses),
- Cash Flow Statement (operating cash flows),
- Financial Ratios (performance and risk measures),

Financial ratios summarize the performance and risk measures of a company, and can be used for investment decisions, like value investing.

```
> # Read table of fundamental data into data frame
> dframe <- read.csv(file="/Users/jerzy/Develop/lecture_slides/data/
```

Balance.Sheet	Income.Statement	Cash.Flow.Statement	Financial.Ratios
Cash & Equivalents	Total Revenue	Cash from Operations	Earnings per Share
Investments	Cost of Goods Sold	Cash from Investments	Price/Earnings Ratio
Receivables	Sales, G and A	Cash from Financing	EBIT
Property & Equipment	Research & Development	FX Adjustment	EBITDA
Goodwill & Intangibles	Depreciation & Amortization	Capital Expenditure	Dividend Yield
Total Assets	Operating Expenses	Starting Cash	Gross Margin
Accounts Payable	Operating Income	Ending Cash	Profit Margin
Long-Term Debt	Investment Income		Free Cash Flow
Total Liabilities	Non-Operating Income		Asset Turnover
Common Stock	Pre-Tax Income		Inventory Turnover
			Receivable Turnover

Sources of Fundamental Company Data

Fundamental company data is obtained from financial statements provided to government agencies and from earnings statements provided to investors.

U.S. companies are required to provide quarterly *10Q* reports and annual *10K* reports to the *Securities and Exchange Commission (SEC)*.

The *10Q* and *10K* reports are publicly available on the *SEC EDGAR* website.

The data provided by *EDGAR* only goes back to the year 2009, because the *SEC* did not mandate electronic filing before then.

The data from the *10Q* and *10K* reports is compiled by vendors such as *Compustat* and *FactSet*, into premium databases.

There are also free databases of fundamental company data, from providers such as *SimFin*, *usfundamentals*, and *Quandl*.

WRDS redistributes fundamental company data from *Compustat*, *S&P Capital IQ*, *Thomson Reuters*, *FactSet*, *Markit*, etc.

There are occasional gaps in the quarterly data, because companies are not required to file quarterly reports on dates when they're also filing annual reports.

The report harmonization procedure fails when companies change their accounting treatment from year to year.

Quandl has written a [review](#) of both premium and free sources of fundamental company data.

Identifiers of Company Data

Identifiers are strings that correspond to *corporate entities* and *securities*.

Headers are current identifiers, while *historical identifiers* have date ranges when they were valid.

Some *universal identifiers* are company namev, tickers, CUSIP, SEDOL, ISIN, and CIK.

Most *universal identifiers*, such as company namev, tickers, CUSIPs, etc. change over time due to corporate events such as mergers, bankruptices, etc.

Tickers are short and recognizable strings, but are not unique, and can change over time and can be reused by a different company.

CUSIPs change infrequently and are never reused, and identify securities in the U.S. and Canada.

SEDOL identify securities in the UK.

ISIN identify international securities in the UK.

CIK identify corporations and individuals who have filed disclosures with the SEC.

Data vendors such as *Compustat*, *Capital IQ*, *IBES*, *Markit*, and *FactSet* have introduced *proprietary identifiers* that are *permanent*.

CRSP *PERMCO* and *PERMNO* are identifiers of companies and securities.

Compustat *GVKEYs* and *IIDs* are identifiers of companies and securities.

The *PERMCO* and *PERMNO* are *permanent identifiers* which never change, even if other identifiers such as tickers do change.

So analysts use *permanent identifiers*, such as *PERMCO*, *PERMNO*, and *GVKEY*.

Capital IQ *Company IDs* are identifiers of companies.

IBES *tickers* are identifiers of securities.

Markit *RED codes* are identifiers of companies.

Factset *Entity IDs* and *Perm Sec IDs* are identifiers of companies and securities.

Wharton Research Data Services WRDS

Wharton Research Data Services (*WRDS*) is a distributor of premium third party data for the academic and research communities.

WRDS provides time series of security prices and fundamental company data, and other financial, econometric, and social datasets.

WRDS provides stock prices, options and implied volatilities, stock fundamentals, financial ratios, zoo::indexes, earnings estimates, analyst ratings, etc.

WRDS redistributes fundamental company data from *Compustat*, *S&P Capital IQ*, *Thomson Reuters*, *FactSet*, *Hedge Fund Research*, *Markit*, etc.

NYU students can obtain user accounts for *WRDS* data.

Subscriptions

- Bank Regulatory
- Beta Suite by WRDS
- Blockholders
- CBOE Indexes
- Compustat - Capital IQ
- Contributed Data
- CRSP
- CSMAR
- DMEF Academic Data
- Dow Jones
- Financial Ratios Suite by WRDS
- IBES
- IHS Global Insight
- Infotrip
- Institutional Shareholder Services (ISS)
- Intraday Indicators by WRDS
- IRI
- Linking Suite by WRDS
- OTC Markets
- Penn World Tables
- Peters and Taylor Total Q
- PHLX
- Public
- Research Quotient
- SEC Order Execution
- TAQ
- Thomson Reuters
- TRACE

Analytics by WRDS

SEC Analytics Suite

Researchers will gain an overview of WRDS' powerful SEC research platform. Some key features they will learn about include:

- Access to 15 million filings from a single index
- Full-text searching over 3.5 million filings
- Sentiment analysis

U.S. Daily Event Study

Investigate abnormal stock returns/volumes around event dates by uploading your own "events" file, or analyzing reaction to firm-specific events from Capital IQ's Key Development database.

- Instant visualization of the effect of events on U.S. equities
- Output includes statistics, plots

Intraday Indicators

Stock specific daily and intraday (5 min, 15 min and 30 min) indicators created from the TAQ intraday datasets.

- Stock specific daily and intraday (5 min, 15 min and 30 min) indicators
- Requires subscription to TAQ dataset.

Classroom by WRDS

WRDS Help and Documentation

WRDS provides extensive *Learning Resources and Documentation*.

WRDS provides online help and a guide to its datasets:

Getting Started

WRDS Data

Finding Data

Understanding Identifiers

Company Identifiers

Futures Contracts

Using SAS Studio

Classroom Tools by wrds

Classroom Tools by WRDS is a teaching and learning toolkit designed specifically for faculty who are introducing finance and business concepts in the classroom.

FEATURED TOOLS:

- Treasury Yield Curve
- Fiscal Policy: Multiplier Effect

Navigation: All Topics | Accounting | Fixed Income | Introduction to WRDS | Investments | Macroeconomics | SEC | Test Analytics

Getting Started

This teaching tool guides your students through their first use of the WRDS website. With it, you can:

- Introduce new users to WRDS
- Provide an overview of how queries are organized
- Demonstrate a sample query to retrieve pricing data

What Data is in WRDS?

Learn about WRDS data offerings and how to navigate them effectively. Request students with the breadth of data sets in WRDS, including:

- Financial
- Economic
- Marketing
- Public

Finding the Data You Want

WRDS contains a tremendous variety of data. This exercise covers the following:

- How data is categorized
- Techniques to locate data
- Resources that make WRDS data


Understanding Identifiers

This exercise describes the major descriptors used to clarify both companies and securities in WRDS, along with examples of each. Students will be introduced to the following:

- Difference between ticker and historical identifiers
- Name vs. company
- Dates of applicability

WRDS Vendor and Database Manuals

WRDS provides many *Vendor and Database Manuals*.



The screenshot shows the WRDS website interface. At the top, there's a navigation bar with the Wharton logo, 'wrds WHARTON RESEARCH DATA SERVICES', a search bar, and links for 'Get Data', 'Analytics', 'Classroom', 'Research', and 'Support'. Below the navigation bar, the page title is 'Manuals and Overviews'. A subtitle reads: 'Manuals, Overviews, and Support Content related to Vendors and Products on the WRDS Platform.' The main content area displays a grid of 30 items, each with a right-pointing arrow and a vendor/product name. The items are arranged in three columns and ten rows.

> AHA	> Global Insight	> Research Quotient
> Audit Analytics	> HFR	> Sustainalytics
> Bank Regulatory	> IBIS/E	> TAQ
> Blockholders	> IMS	> TRACE
> BoardEx	> IRI	> Thomson Reuters
> Bureau van Dijk	> ISS (Formerly RiskMetrics)	> Total_c
> CIBC Indexes	> ISSM	> Toyo Keizai
> CENTRIS	> Infogroup	> TwoIQ
> CRSP	> KLD	> WRDS Bag of Words
> CSMAR	> Kdmine	> WRDS Factors (Backtester)
> CUSIP	> LSPD	> WRDS Beta Suite
> Clarivate	> Levin	> WRDS Bond Return
> Compustat	> MFLINKS	> WRDS Efficient Frontier
> Compore	> MSGI	> WRDS Event Study Suite
> DAMEF	> MSTRB	> WRDS Financial Ratios

Downloading Data From WRDS

The two main databases on WRDS are the *Compustat* database of fundamental company data and security prices, and the *CRSP* database of security prices.

The WRDS data can be accessed through vendor pages, which are conveniently organized by *concept*.

The easiest way to download data from WRDS into R, is by first downloading it into a .csv file, and then reading it into R.

The screenshot shows the Wharton WRDS website interface. At the top, there's a navigation bar with the Wharton logo and 'wrds' branding. Below the navigation bar, the main heading is 'Browse Data by Concept'. On the right side, there's a 'Concept Finder' search box. The main content area lists several categories of data available for download:

- Company Financials**
 - Financial Statements
 - SAP Compustat
 - Thomson Reuters Worldscope
 - Factset Fundamentals
 - BVD Analytics
 - BVD Orlis
 - PACAP Pacific Basin
 - CSMAR Financials - China Stock Market & Accounting Research
 - SAP Capital IQ Capital Structure
 - Gaibench As-Reported Filings & Footnotes
- Audit & Regulatory Filings**
 - WRDS SEC Analytics Suite
 - Audit Analytics
 - SAP Filings Database
- Banks**
 - SAP Compustat Bank
 - SNL
 - BVD Orlis Bank Focus
 - Bank Regulatory Call Reports - Federal Reserve Bank of Chicago
 - FRB H15 Interest Rates
- Segments / Industry Data**
 - SAP Compustat Segments
 - Thomson Reuters Worldscope Segments
 - SAP Compustat Industry-Specific
 - Thomson Reuters Ibs RPI
 - Factset Revers
 - BVD Iis Insurance Companies
- Compensation**
 - SAP Compustat Execucomp
 - SAP Capital IQ People Intelligence
 - ISS Incentive Lab
 - BoardEx
 - MSCI: GMI Ratings
- Intellectual Property**
 - USPINE
- Financial Markets, Prices, Returns**
 - Stock Prices
 - CRSP
 - Compustat
 - Factset

At the bottom right of the page, there is a 'Top of Section' link.

The Compustat Database on WRDS

Compustat is a provider of fundamental company data, and also security prices, and other data such as credit ratings.

Daily North American security prices can be downloaded from the *Compustat Daily Prices* database, which can be reached by navigating from the *WRDS* home page to *North America - Daily* and then to *Security Daily*.

Compustat is part of *S&P Capital IQ*.

Compustat uses the *GVKEY* and *IID* identifiers of companies and securities.

Compustat provides only recent time series of data starting in 1983.

The top screenshot shows the Wharton WRDS Compustat - Capital IQ from Standard & Poor's page. It includes a search bar, navigation links, and a list of databases. The databases are categorized into Compustat, Capital IQ, and Other Compustat. The bottom screenshot shows the North America - Daily page, which provides a detailed description of the database and lists various data series.

Compustat - Capital IQ from Standard & Poor's

Learn more about SAP Global Market Intelligence data through WRDS >>

For more about this dataset, see the [Dataset List](#), [Manuals and Overviews](#) or [FAQs](#).

Compustat	Capital IQ	Other Compustat
Databases in this section are updated every day unless otherwise noted. Update schedules should not be confused with end-of-day or end-of-month data such as stock prices.	Capital IQ is a suite of databases from SAP. They connect to the Compustat family of databases through GVKEY.	
> North America - Daily (13) > Global - Daily (5) > Bank - Daily (2) > Historical Segments - Daily (2) > Execucomp - Monthly Updates (2)	> Capital Structure (3) > Identifiers (4) > Key Developments (1) > People Intelligence (3) > Transactions (2) > Transcripts (3)	> North America - Annual Updates (13) > Marginal Tax Rates (1) > SAP Filing Dates (1) > Preliminary History (2) > Unrestated Quarterly (2) > Point in Time (2) > North America - Daily Updates (Non-Historical) (13)

North America - Daily

For more about this dataset, see the [Dataset List](#), [Manuals and Overviews](#) or [FAQs](#).

Compustat North America is a database of U.S. and Canadian fundamental and market information on active and inactive publicly held companies. It provides more than 300 annual and 100 quarterly Income Statement, Balance Sheet, Statement of Cash Flows, and supplemental data items. For most companies, annual history is available back to 1950 and quarterly history back to 1962 with monthly market history back to 1962.

Compustat North America files also contain information on aggregates, industry segments, banks, market prices, dividends, and earnings.

Fundamentals Annual	Industry Specific Quarterly	Segments (Non-Historical)
Fundamentals Quarterly	Pension Annual	Segments (Non-Historical) - Customer
Index Constituents	Pension Quarterly	Simplified Financial Statement Extract
Index Fundamentals	Ratings	Supplemental Short Interest File
Index Prices	Security Daily	Financial Ratios Industry Level
Industry Specific Annual	Security Monthly	Financial Ratios Firm Level

Package *rwrds* for Downloading Data From *WRDS*

The package *rwrds* contains functions for downloading data from *WRDS*.

```
> # Install package "rwrds"
> devtools::install_github("davidsovlch/rwrds")
> # Load package "rwrds"
> library(rwrds)
> # Get documentation for package "rwrds"
> packageDescription("rwrds")
> # Load help page
> help(package="rwrds")
> # List all datasets in "rwrds"
> data(package="rwrds")
> # List all objects in "rwrds"
> ls("package:rwrds")
> # Remove rwrds from search path
> detach("package:rwrds")
```

Downloading Names Table From *Compustat*

The *Compustat* names table is a database containing various company *identifiers* (company namev, tickers, CUSIPs, GVKEYs).

The names table can be used to cross-reference the *identifiers*, for example to find the *GVKEYs* from the company *tickers*.

The function `rwrds::compustat_names()` downloads the names table from *Compustat*.

```
> library(rwrds)
> library(dplyr)
> # Establish connection to WRDS
> connv <- rwrds::connvnect(username="jp3900", password="RipvanWinkl
> # Download Compustat names table as dplyr object
> namet <- rwrds::compustat_names(wrds=connv, subset=FALSE, dl=TRUE)
> dim(namet)
> # Save names table as csv file
> write.csv(namet, file="/Users/jerzy/Develop/lecture_slides/data/co
> # rm(namet)
> # Read names table from csv file
> namet <- read.csv(file="/Users/jerzy/Develop/lecture_slides/data/
> # symboln <- "VTI"
> # match(symboln, namet$tic)
> # Create ETF symbols (tickers)
> symbolv <- c("VTI", "VEU", "EEM")
> # Get cusips of symbolv
> indeks <- match(symbolv, namet$tic)
> names(indeks) <- symbolv
> etf_cusips <- namet$cusip[indeks]
> names(etf_cusips) <- symbolv
> # Save cusips into text file
> cat(etf_cusips, file="/Users/jerzy/Develop/lecture_slides/data/etf
> # Save gvkeys into text file
> etf_gvkeys <- namet$gvkey[indeks]
> names(etf_gvkeys) <- symbolv
> cat(etf_gvkeys, file="/Users/jerzy/Develop/lecture_slides/data/etf
```

draft: Downloading Index Constituents From *Compustat*

The names and tickers of the constituents of a given index can be downloaded from the *WRDS Compustat web page*.

The constituents can be downloaded for a specific date, or for a range of dates, to obtain the names and tickers of all the companies which belonged to the index during that time.

The *S&P500 Index* ticker is equal to *i0003*, and its *GVKEY* is equal to 3.

WRDS provides *help about how to download index constituents*. The table with the index constituents can be read into R a .csv from the file downloaded from the *WRDS Compustat web page*.

The table of constituents is a data frame containing various company *identifiers* (company namev, tickers, CUSIPs, GVKEYs).

Some *GVKEYs* are duplicates because some companies entered the index more than once.

The *GVKEYs* and *CUSIPs* can be saved into text files for use in *WRDS* data queries.

```
> # Read .csv file with S&P500 constituents
> sp500table <- read.csv(file="/Users/jerzy/Develop/lecture_slides/")
> class(sp500table)
> dim(sp500table)
> head(sp500table)
> # Select unique sp500 tickers and save them into text file
> ticksp500 <- unique(sp500table$co_tic)
```

The screenshot shows the WRDS website interface. On the left is a navigation menu with options like 'Fundamentals Annual', 'Index Constituents', etc. The main content area is titled 'Compustat - Capital IQ' and 'North America - Daily'. It contains a 'Query Form' section with two steps: 'Step 1: Choose your date range' and 'Step 2: Apply your company codes'. Step 1 includes a date range selector (2010-01 to 2020-07). Step 2 includes a company code selector where 'i0003' is entered. Below this is a 'Query Variables' section with a search bar and a list of variables including 'CDIND - Index Name', 'INDLACAT - Index Category Code', 'INDLTYPE - Index Code Type', 'SIF1 - S&P Industry Index Identifier', 'SIFM - S&P Major Index Identifier', 'TIC - Ticker', 'CDIND - Company Name', 'GVKEY - GVKEY', and 'CUSIP - Cusip'.

Downloading the S&P500 Index Constituents

As of July 2020, the *S&P500* stock index constituents data was removed from *Compustat*, so it's no longer available on *WRDS*.

As an alternative, users can download the *SPY ETF* constituents from *State Street*.

The file `sp500_constituents.csv` contains the symbols (tickers) for almost 800 stocks belonging to the *S&P500* stock index, either now or in the past.

```
> # Read .csv file with S&P500 constituents
> sp500table <- read.csv(file="/Users/jerzy/Develop/lecture_slides/c
> class(sp500table)
> dim(sp500table)
> head(sp500table)
```

Downloading Fundamental Company Data

Fundamental company data can be downloaded from the *Compustat Fundamentals* database.

The user can download data for either a single identifier, or for many identifiers supplied in a text file.

Wharton wrds WHARTON RESEARCH DATA SERVICES

Home | Get Data | Compustat | Capital IQ | Compustat | North America - Daily | Compustat Daily Updates - Fundamentals Quarterly

Query Form | Variable Descriptions | Manuals and Overview | FAQs | Contact List

Compustat Daily Updates - Fundamentals Quarterly

Step 1: Choose your date range.

Date Variable:

Date range: 2015-12 to 2020-07

Step 2: Apply your company codes.

☒ IIC ☐ GVKEY ☐ CUSIP ☐ SIC ☐ NAICS ☐ CIK ☐ GIC Sub-Industry

Select an option for entering company codes

☒ MSF ☐ Code List Name

Please enter Company codes separated by a space.
Example: IBM MSFT DELL / Code Lookup

No file selected

Upload a plain text file (.txt), having one code per line.

☐ Select Saved CodeList

Choose from your saved code lists.

☐ Search the entire database

This method allows you to search the entire database of records. Please be aware that this method can take a very long time to run because it is dependent upon the size of the database.

Screening Variables

Several screening variables are pre-selected to produce one record per GVKEY-DATADATE pair, while keeping the vast majority of records. Examples of excluded rows include those with restated data, different views of the same data (pre-form, pre-PAGE). Click on each variable for a more detailed explanation.

Consolidation Level	SIC	INDL	INDS	INDS	Output
Industry Format	☐	☐	☐	☐	☐
Date Format	☒	☐	☐	☐	☐
Population Source	☒	☐	☐	☐	☐
Quarter Type	☐	☐	☐	☐	☐
Currency	☐	☐	☐	☐	☐
Company Status	☐	☐	☐	☐	☐

Step 3: Choose variable types.

Select Variable Types

Select the items you would like to include in your search.

Select (2) Selected (0) Clear All (1)

Downloading Daily Prices From Compustat

Daily security prices can be downloaded from the *Compustat Daily Prices* database.

The user can download data for either a single identifier, or for many identifiers in a text file.

The screenshot shows the Wharton WRDS interface. The left sidebar lists various data sources, with 'Security Daily' selected. The main panel is titled 'Compustat Daily Updates - Security Daily'. It contains a 'Query Form' tab and a 'Variable Descriptions' tab. The 'Query Form' tab has a 'Date range' section with a date picker set to 2010-01-01 to 2020-01-01. Below this is a 'Step 2: Apply your company codes.' section with a 'Select an option for entering company codes.' dropdown menu. The 'Security Daily' option is selected. The 'Step 3: Query Variables.' section is also visible, showing a list of variables to be included in the query.

The query for *OHLC* prices should also include the *split adjustment factor* and the *total return factor*.

Step 3: Query Variables.

How does this work?

The screenshot shows the 'Identifying Information' tab of the Wharton WRDS interface. It displays a list of variables to be included in the query. The 'Select' column has a dropdown menu set to 'All'. The 'Selected' column has a dropdown menu set to 'Clear All'. The list of variables includes: DIVSPAYDATE - Special Cash Dividends - Daily Payment Date, DVI - Indicated Annual Dividend - Current, DIVRATED - Indicated Annual Dividend Rate - Daily, EPS - Current EPS, EPSMO - Current EPS Month, IID - Issue ID - Dividends, PAYDATE - Dividend Payment Date, PAYDATEIND - Dividend Payment Date Indicator, PRICSD - Price Status Code - Daily, and RECORDDATE - Dividend Record Date. The 'Selected' column shows the following variables: Company Name, Ticker Symbol, CUSIP, PRCOD - Price - Open - Daily, PRCHD - Price - High - Daily, PRCLD - Price - Low - Daily, PRCCD - Price - Close - Daily, CSHTRD - Trading Volume - Daily, AJXEDI - Adjustment Factor (Issue)-Cumulative by Ex-Date, and TRFD - Daily Total Return Factor.

Reading Daily Prices Into R

The time series data downloaded into a .csv file can then be read into R.

But the data downloaded from the *Compustat Daily Prices* database may contain data for several related securities, so it must be properly selected (pruned).

The *OHLC* prices can be adjusted by dividing them by the *split adjustment factor* and multiplying them by the *total return factor*.

The *OHLC* prices are also divided by the final *total return factor* so that the most recent adjusted prices are equal to the most recent market prices.

```
> # Read .csv file with TAP OHLC prices
> ohlc <- read.csv(file="/Users/jerzy/Develop/lecture_slides/data/TAP.csv")
> # ohlc contains cusips not in cusisps500
> cusips <- unique(ohlc$cusip)
> cusips %in% cusisps500
> # Select data only for cusisps500
> ohlc <- ohlc[ohlc$cusip %in% cusisps500, ]
> # ohlc contains tickers not in ticksp500
> tickers <- unique(ohlc$tic)
> tickers %in% ticksp500
> # Select data only for ticksp500
> ohlc <- ohlc[ohlc$tic %in% ticksp500, ]
> # Select ticker from sp500table
> symboln <- sp500table$co_tic[match(ohlc$gvkey[1], sp500table$gvkey)]
> # Adjustment factor and total return factor
> factadj <- drop(ohlc[, "ajexdi"])
> factr <- drop(ohlc[, "trfd"])
> # Extract index of dates
> datev <- drop(ohlc[, "datadate"])
> datev <- lubridate::ymd(datev)
> # Select only OHLCV data columns
> ohlc <- ohlc[, c("prcd", "prchd", "prcld", "prccd", "cshtd")]
> colnames(ohlc) <- paste(symboln, c("Open", "High", "Low", "Close", "Volume"))
> # Coerce to xts series
> ohlc <- xts::xts(ohlc, datev)
> # Fill the missing (NA) Open prices
> isna <- is.na(ohlc[, 1])
> ohlc[isna, 1] <- (ohlc[isna, 2] + ohlc[isna, 3])/2
> sum(is.na(ohlc))
> # Adjust all the prices
> ohlc[, 1:4] <- factr*ohlc[, 1:4]/factadj/factr[NROW(factr)]
> ohlc <- na.omit(ohlc)
> plot(quantmod::Cl(ohlc), main="TAP Stock")
```

Function for Reading Daily Prices Into R

The function `format_ohlc()` formats raw data downloaded from *Compustat* into a time series of *OHLC* prices.

```
> # Define formatting function for OHLC prices
> formohlc <- function(ohlc, pricenv) {
+   ## Select ticker from sp500table
+   symboln <- sp500table$co_tic[match(ohlc$gvkey[1], sp500table$gvkey)]
+   ## Split adjustment and total return factors
+   factadj <- drop(ohlc[, c("ajexdi")])
+   factr <- drop(ohlc[, "trfd"])
+   factr <- ifelse(is.na(factr), 1, factr)
+   ## Extract dates index
+   datev <- drop(ohlc[, "datadate"])
+   datev <- lubridate::ymd(datev)
+   ## Select only OHLCV data
+   ohlc <- ohlc[, c("prcod", "prchd", "prcld", "prccd", "cshtd")]
+   colnames(ohlc) <- paste(symboln, c("Open", "High", "Low", "Close", "Volume"))
+   ## Coerce to xts series
+   ohlc <- xts::xts(ohlc, datev)
+   ## Fill NA prices
+   isna <- is.na(ohlc[, 1])
+   ohlc[isna, 1] <- (ohlc[isna, 2] + ohlc[isna, 3])/2
+   ## Adjust the prices
+   ohlc[, 1:4] <- factr*ohlc[, 1:4]/factadj/factr[NROW(factr)]
+   ## Copy the OHLCV data to pricenv
+   ohlc <- na.omit(ohlc)
+   assign(x=symboln, value=ohlc, envir=pricenv)
+   symbol
+ } ## end formohlc
```


Reading Index Constituent Prices Into R

The prices for all the index constituents are first downloaded from *Compustat* into a single .csv file, and then they are read into R and formatted into a time series of *OHLC* prices.

But the prices downloaded from the *Compustat Daily Prices* database often contain prices for several related securities from the same company, so they must be properly selected (pruned) to match the index constituents.

```
> # Load OHLC prices from .csv file downloaded from WRDS by cusip
> pricesp500 <- read.csv(file="/Users/jerzy/Develop/lecture_slides/compustat_prices.csv")
> # sp500_prices contains cusips not in cusipsp500
> cusips <- unique(sp500_prices$cusip)
> NROW(cusipsp500); NROW(cusips)
> # Select data only for cusipsp500
> pricesp500 <- pricesp500[pricesp500$cusip %in% cusipsp500, ]
> # sp500_prices contains tickers not in ticksp500
> tickers <- unique(sp500_prices$tic)
> NROW(ticksp500); NROW(tickers)
> # Select data only for ticksp500
> pricesp500 <- pricesp500[pricesp500$tic %in% ticksp500, ]
> # Create new data environment
> sp500env <- new.env()
> # Perform OHLC aggregations by cusip column
> pricesp500 <- split(pricesp500, sp500_prices$cusip)
> process_ed <- lapply(pricesp500, formohl, pricenv=sp500env)
```

Managing Exceptions in Index Constituent Prices

Some securities of past index constituents may no longer be trading because of corporate events (mergers, bankruptcies, etc.), so they should be removed from the data.

Tickers which contain a dot in their name (like "BRK.B") are not valid symbols in R, so they must be renamed.

The column names for symbol "LOW" (Lowe's company) must be renamed for the extractor function `quantmod::Lo()` to work properly.

```
> # Get end dates of series in sp500env
> endd <- eapply(sp500env, end)
> endd <- unlist(endd)
> endd <- as.Date(endd)
> # Remove elements with short end dates
> ishort <- (endd < max(endd))
> rm(list=names(sp500env)[ishort], envir=sp500env)
> # Rename element "BRK.B" to "BRKB"
> sp500env$BRKB <- sp500env$BRK.B
> rm(BRK.B, envir=sp500env)
> names(sp500env$BRKB) <- paste("BRKB",
+   c("Open", "High", "Low", "Close", "Volume"), sep=".")
> # Rename element "LOW" to "LOWES"
> sp500env$LOWES <- sp500env$LOW
> rm(LOW, envir=sp500env)
> names(sp500env$LOWES) <- paste("LOWES",
+   c("Open", "High", "Low", "Close", "Volume"), sep=".")
> # Rename element "BF.B" to "BFB"
> sp500env$BFB <- sp500env$BF.B
> rm(BF.B, envir=sp500env)
> names(sp500env$BFB) <- paste("BFB",
+   c("Open", "High", "Low", "Close", "Volume"), sep=".")
> # Save OHLC prices to .RData file
> save(sp500env, file="/Users/jerzy/Develop/lecture_slides/data/sp500.RData")
> plot(quantmod::Cl(sp500env$MSFT))
```

Saving WRDS Queries

WRDS queries can be saved and reused again.

The screenshot displays the Wharton WRDS (Worldwide Research Data Services) interface. The main navigation bar at the top includes links for 'Get Data', 'Analytics', and 'Classrooms'. The left sidebar lists various data categories under 'Compustat - Capital IQ', with 'Security Daily' currently selected. The main content area shows the 'Compustat Daily Updates - Security Daily' query form. It includes a 'Date range' section with a date picker set to 2010-01-01 to 2020-07-05. Below this is a 'Step 2: Apply your company codes' section with radio buttons for 'Company Codes', 'Code List Name', and 'Search the entire database'. The 'Company Codes' option is selected. At the bottom, there is a 'Step 3: Query Variables' section with a search bar and a list of variables.

draft: Downloading Fundamental Company Data From *Quandl*

SEC is a free database of stock fundamentals extracted from SEC and (but not harmonized),

<https://www.quandl.com/data/SEC>

RAYMOND is a free database of harmonized stock fundamentals, based on the *SEC* database,

<https://www.quandl.com/data/RAYMOND>

<https://www.quandl.com/data/RAYMOND?keyword=aapl>

Wharton Research Data Services (*WRDS*) is a distributor of third party data for the academic and research communities.

WRDS provides time series of stock prices and fundamental company data, and other financial, econometric, and social datasets.

WRDS redistributes fundamental company data from *Compustat*, *S&P Capital IQ*, *Thomson Reuters*, *FactSet*, *Hedge Fund Research*, *Markit*, etc.

The Center for Research in Security Prices (*CRSP* is a provider of historical security prices for the academic and research communities.

Much of the *WRDS* data is free, while premium data can be obtained under a temporary license.

WRDS provides online help and a guide to its datasets:

<https://wrds-www.wharton.upenn.edu>

WRDS provides stock prices, stock fundamentals, financial ratios, zoo::indexes, options and volatility,

```
> library(Quandl) ## Load package Quandl
> # Register Quandl API key
> Quandl.api_key("pVJi9Nv3V8CD3Js5s7Qx")
>
> # Quandl stock market data
> # https://blog.quandl.com/stock-market-data-ultimate-guide-part-1
> # https://blog.quandl.com/stock-market-data-the-ultimate-guide-part-2
>
> # Download RAYMOND metadata
> # https://www.quandl.com/data/RAYMOND-Raymond/documentation/metadata
>
> # Download S&P500 Index constituents
> # https://s3.amazonaws.com/static.quandl.com/tickers/SP500.csv
>
> # Download AAPL gross profits from RAYMOND
> profitaapl <- Quandl("RAYMOND/AAPL_GROSS_PROFIT_Q", type="xts")
> chart_Series(profitaapl, name="AAPL gross profits")
>
> # Download multiple time series
> pricev <- Quandl(code=c("NSE/OIL", "WIKI/AAPL"),
+   startd="2013-01-01", type="xts")
>
> # Download datasets for AAPL
> # https://www.quandl.com/api/v3/datasets/WIKI/AAPL.json
>
> # Download metadata for AAPL
> pricev <- Quandl(code=c("NSE/OIL", "WIKI/AAPL"),
+   startd="2013-01-01", type="xts")
> # https://www.quandl.com/api/v3/datasets/WIKI/AAPL/metadata.json
>
> # Scrape fundamental data from Google using quantmod - doesn't work
> fundata <- getFinancials("HPQ", src="google", auto.assign=FALSE)
> # View quarterly fundamentals
> viewFinancials(fundata, period="Q")
> viewFinancials(fundata)
>
> # Scrape fundamental data from Yahoo using quantmod
> fundata <- getFinancials("HPQ", src="yahoo", auto.assign=FALSE)
```

The CRSP Database on WRDS

The Center for Research in Security Prices (*CRSP*) is a provider of historical security prices, and is an affiliate of the University of Chicago Booth School of Business.

CRSP data provides daily prices, but it's updated only monthly, quarterly, and annually.

CRSP provides very long time series of data starting in 1926.

CRSP data is indexed using the *PERMCO* company identifier and the *PERMNO* security identifier.

Since a single company can have many securities, therefore a single *PERMCO* may be linked to multiple *PERMNO*s.

WRDS provides a tool to translate tickers to *PERMCO*/*PERMNO*:

https://wrds-web.wharton.upenn.edu/wrds/ds/crsp/tools_a/dse/translate/index.cfm

The screenshot shows the Wharton WRDS website interface. At the top, there's a navigation bar with 'Wharton wrds' logo and a search bar. Below the navigation bar, the main heading is 'The Center for Research in Security Prices (CRSP)'. A note states: 'For more about this dataset, see the Dataset List, Manuals and Overviews or FAQs.' Below this, there are sections for 'Annual Update', 'Quarterly Update', and 'Monthly Update'. Each section contains a list of database categories with their respective update frequencies. For example, under 'Annual Update', categories include 'Stock / Security Files' (2), 'Index / Treasury and Inflation' (1), 'Tools' (5), 'Stock / Events' (3), 'Index / CRSP Select Series' (5), 'Stock-1962 / Security Files' (4), 'Stock / Portfolio Assignments' (5), 'CRSP/Computat Merged' (5), 'Stock-1962 / Events' (5), 'Index / Stock File Indexes' (5), 'Treasures' (11), 'Ziman REIT' (4), 'Index / Cap-Based Portfolios' (2), 'Treasury / Daily (Legacy)' (5), and '10 Year U.S. Stock' (1). Similar lists are provided for 'Quarterly Update' and 'Monthly Update'.

Downloading Daily Prices From CRSP

Daily North American security prices can be downloaded from the *CRSP daily* database.

The time series data can be downloaded into a .csv file, and then read into R.

```
> # Specify class for column "TICKER" so that "F" doesn't become FALSE
> colv <- "character"
> names(colv) <- "TICKER"
> # Read .csv file with Ford OHLC prices
> ohlc <- read.csv(file="/Users/jerzy/Develop/lecture_slides/data/F",
+   colClasses=colv)
> symboln <- ohlc[, "TICKER"]
> # Adjustment factor
> factadj <- drop(ohlc[, "CFACPR"])
> # Extract dates index
> datev <- drop(ohlc[, "date"])
> datev <- lubridate::ymd(datev)
> # Select only OHLCV data
> ohlc <- ohlc[, c("OPENPRC", "ASKHI", "BIDLO", "PRC", "VOL")]
> colnames(ohlc) <- paste(symboln, c("Open", "High", "Low", "Close"))
> # Coerce to xts series
> ohlc <- xts::xts(ohlc, datev)
> # Fill missing Open NA prices
> isna <- is.na(ohlc[, 1])
> ohlc[isna, 1] <- (ohlc[isna, 2] + ohlc[isna, 3])/2
> # Adjust all the prices
> ohlc[, 1:4] <- ohlc[, 1:4]/factadj/factr[NROW(factr)]
> plot(quantmod::Cl(ohlc), main="Ford Stock")
```

The screenshot displays the Wharton Research Data Services (WRDS) interface for the CRSP Daily Stock database. The top navigation bar includes links for Home, Get Data, Annual Update, Stock Security Data, CRSP Daily Stock, and other resources. The main content area is divided into several sections:

- CRSP**: A sidebar menu with options like Monthly Stock File, Daily Stock File, Stock Market Indices, Stock Header Info, Beta Suite by WRDS, U.S. Daily Event Study, and Linking Tools.
- CRSP Daily Stock**: The main section, which includes a 'Query Form' tab and a 'Variable Descriptions' tab.
- Step 1: Choose your date range.**: A date range selector showing 1926-01-01 to 2019-12-31.
- Step 2: Apply your company codes.**: A section for selecting company codes, with options for 'F' (Ford) and 'C' (Comp).
- Step 3: Query Variables.**: A section for selecting variables, with a list of variables including Comp, Numup, Bid, and others.

draft: The Merged *CRSP/Compustat* Database

Links:

<http://kaichen.work/?p=138>

<http://www.crsp.org/products/documentation/crsp-link>

<https://mingze-gao.com/posts/merge-compustat-and-crsp/>

<http://www.crsp.org/products/documentation/crspcompustat-merged-database-guide>

The Center for Research in Security Prices (*CRSP* is a provider of historical security prices.

CRSP data is indexed using the *PERMCO* company identifier and the *PERMNO* security identifier.

Therefore a given *PERMCO* may be linked to multiple *PERMNOs*.

Compustat provides financial statement data of a company. The micro unit on Compustat is each and every company. However, it is less known that Compustat also provides security data. In addition, because the coverage of Compustat is more extensive than that of *CRSP*, Compustat contains additional security data that are unavailable on *CRSP*.

Compustat uses *IID* and *GVKEY* to identify all securities. A marker item, *PRIMISS*, indicates whether a security is primary or secondary.

Similar to *PERMCO* on *CRSP*, one *GVKEY* may have multiple *IIDs* assigned to it.

Database	Identifier	Description
CRSP	PERMCO	Permanent company identifier.
CRSP	PERMNO	Permanent security identifier. One PERMNO belongs to only one PERMCO. One PERMCO may have one or more PERMNOs.
Compustat	GVKEY	Permanent company identifier.
Compustat	IID	Permanent security identifier. One GVKEY may have multiple IIDs. Both are needed to identify a Compustat security.
Compustat	PRIMISS	Indicator for primary (P) or secondary (J) securities.

draft: Downloading Daily Prices From Merged *CRSP/Compustat* Database

Daily North American security prices can be downloaded from the *CRSP/Compustat daily* database.

The time series data can be downloaded into a .csv file, and then read into R.

```
> # Read .csv file with TAP OHLC prices
> ohlc <- read.csv(file="Users/jerzy/Develop/lecture_slides/data/TAP_2015.csv")
> symboln <- ohlc[1, "tic"]
> # Adjustment factor and total return factor
> factadj <- drop(ohlc[, "ajexdi"])
> factr <- drop(ohlc[, "trfd"])
> # Extract dates index
> datev <- drop(ohlc[, "datadate"])
> datev <- lubridate::ymd(datev)
> # Select only OHLCV data
> ohlc <- ohlc[, c("prcdd", "prchd", "prclld", "prccd", "cshtd")]
> colnames(ohlc) <- paste(symboln, c("Open", "High", "Low", "Close", "Volume"))
> # Coerce to xts series
> ohlc <- xts::xts(ohlc, datev)
> # Fill missing Open NA prices
> isna <- is.na(ohlc[, 1])
> ohlc[isna, 1] <- (ohlc[isna, 2] + ohlc[isna, 3])/2
> sum(is.na(ohlc))
> # Adjust all the prices
> ohlc[, 1:4] <- factr*ohlc[, 1:4]/factadj[factr[NROW(factr)]]
> plot(quantmod::CL(ohlc), main="TAP Stock")
```



UNIVERSITY OF PENNSYLVANIA

wharton

FINANCIAL SERVICES

[Get Data](#)
[Analytics](#)
[Classroom](#)
[Research](#)
[Support](#)

[JEREMY'S ACCOUNT](#)
[LOGOUT](#)

[Home](#)
[Get Data](#)
[Computat - Capital IQ](#)
[Computat](#)
[North America - Daily](#)
[Computat Daily Updates - Security Daily](#)

[Search WFINQ](#)

Computat - Capital IQ

North America - Daily

Fundamentals Annual

Fundamentals Quarterly

Index Constituents

Index Fundamentals

Index Prices

Industry Specific Annual

Industry Specific Quarterly

Pension Annual

Pension Quarterly

Ratings

Security Daily

Security Monthly

Segments (Non-Historical)

Segments (Non-Historical) - Customer

Simplified Financial Statement Extract

Supplemental Short Interest File

Financial Ratios Industry Level

Financial Ratios Firm Level

Computat Daily Updates - Security Daily

You have 1 saved queries for this dataset.

Step 1: Choose your date range.

Date Variables: Datefields

Date range: to

Step 2: Apply your company codes.

☐ SEC ☐ GWSLY ☐ CUSIP ☐ SIC ☐ NAICS ☐ OK ☐ SIC Sub-Industry

Select an option for entering company codes

☒ MSF

Enter entire Company codes separated by a space.
Example: IBM /MSFT /GGL / Code Lookup /

☐ Code List Name

Press code list to Saved Codes

Upload a file (not the .csv), having one code per line.

☐ -----Select Saved Codes-----

Choose from your saved code lists

☐ Choose the entire database

The method allows you to search the entire database of records. Please be aware that this method can take a very long time to run because it is dependent upon the size of the database.

Step 3: Query Variables.

How does this work?

[View](#)

Select (24)

- ☒ DIVIDENDPAYOUT - Cash Dividends - Daily Payment Date Indicator
- ☒ DIVSP - Special Cash Dividends - Daily
- ☒ DIVSPRINCALS - Special Cash Dividends - Daily Payment Date
- ☒ DIVI - Initiated Annual Dividend - Current
- ☒ DIVINITIATD - Initiated Annual Dividend - Future Daily
- ☒ EPS - Current EPS
- ☒ EPSMAD - Current EPS Month
- ☒ RD - Basic ID - Dividends
- ☒ PAYDATE - Dividend Payment Date

[View](#)

Selected (10)

- ☒ Index Symbol
- ☒ CUSIP
- ☒ ADJUSTC - Adjustment Factor (Stocks) - Continuity by Cn-Date
- ☒ CSRTIHD - Trading Volume - Daily
- ☒ PRECLO - Price - Close - Daily
- ☒ PREHLS - Price - Low - Daily
- ☒ PREHGH - Price - High - Daily
- ☒ PREHLS - Price - Low - Daily
- ☒ PREHGH - Price - High - Daily
- ☒ TRFD - Daily Total Return Factor

Downloading Fama-French Factors from *Quandl*

The Fama/French factors are constructed using six value-weight portfolios formed on size and book-to-market,

<https://www.quandl.com/data/KFRENCH/FACTORS.D>

Mkt-RF is the excess return on the market (value-weighted NYSE, AMEX, and NASDAQ stocks minus the one-month Treasury bill rate).

SMB (Small Minus Big) is the return on the three small-cap portfolios minus the three big-cap portfolios.

HML (High Minus Low) is the return on the two value portfolios minus the two growth portfolios.

```
> # Download Fama-French factors from KFRENCH database
> ffactors <- Quandl(code="KFRENCH/FACTORS_D",
+   startd="2001-01-01", type="xts")
> dim(ffactors)
> head(ffactors)
> tail(ffactors)
> chart_Series(cumsum(ffactors["2001/", 1]/100),
+   name="Fama-French factors")
```

Trade and Quote (TAQ) Data

High frequency data is typically formatted as either Trade and Quote (TAQ) data, or *Open-High-Low-Close* (OHLC) data.

Trade and Quote (TAQ) data contains intraday *trades* and *quotes* on exchange-traded stocks and futures.

TAQ data is often called *tick data*, with a *tick* being a row of data containing new *trades* or *quotes*.

The TAQ data is spaced irregularly in time, with data recorded each time a new trade or quote arrives.

Each row of TAQ data may contain the quote and trade prices, and the corresponding quote size or trade volume: *Bid.Price*, *Bid.Size*, *Ask.Price*, *Ask.Size*, *Trade.Price*, *Volume*.

TAQ data is often split into *trade* data and *quote* data.

```
> # Load package HighFreq
> library(HighFreq)
> # Or load the high frequency data file directly:
> # symbolv <- load("/Users/jerzy/Develop/R/HighFreq/data/hf_data.R")
> head(HighFreq::SPY_TAQ)
> head(HighFreq::SPY)
> tail(HighFreq::SPY)
```

Downloading TAQ Data From WRDS

TAQ data can be downloaded from the *WRDS TAQ* web page.

The TAQ data are at millisecond frequency, and are *consolidated* (combined) from the New York Stock Exchange *NYSE* and other exchanges.

The *WRDS TAQ* web page provides separately *trades* data and separately *quotes* data.

Wharton RESEARCH DATA SERVICES

Search WRDS

Home / Get Data / TAQ / Millisecond Trade and Quote / TAQ - Millisecond Consolidated Trades

TAQ

Millisecond Trade and Quote

Consolidated Quotes

Consolidated Trades

TAQ Millisecond Tools

Query Form Variable Descriptions Manuals and Overviews FAQs Dataset List

TAQ - Millisecond Consolidated Trades

Note to TAQ users: Missing trades in the April 5, 2012 Consolidated Trades dataset.

There are roughly 65,000 missing trade records for tickers with initial letters M through R from 08:42 through 10:02:50 on April 5, 2012. NYSE has reported to WRDS that NASDAQ was not disseminating trade reports during that time period. The actual number of trade records (for "M" through "R") during that time from NASDAQ (exchange code "Q") was 2,232, but the average during the rest of the week was 65,274. The missing trades amount to about 12% of total trades in those equities.

WRDS will provide an update as soon as the files are available.

You have 1 saved query for this dataset.

Step 1: Choose your date range.
I would like data from Mar 18 2020 to Mar 18 2020

Step 2: What Time Range
Filter observations by time range

Beginning 00 30 00 Ending 18 00 00

Step 3: Apply your company codes.
@ DRMSOL (past partner)

Select an option for entering company codes

☒ AAPL ☐ Code List Name

Please enter Company codes separated by a space.
Example: WMT MFT DELL COB LOOKUP

☐ No Name ☐ No file selected

Reading TAQ Data From .csv Files

Trade and Quote (TAQ) data stored in .csv files can be very large, so it's better to read it using the function `data.table::fread()` which is much faster than the function `read.csv()`.

Each *trade* or *quote* contributes a *tick* (row) of data, and the number of ticks can be very large (hundred of thousands per day, or more).

The function `strptime()` coerces character strings representing the date and time into POSIXlt *date-time* objects.

The argument `format="%H:%M:%OS"` allows the parsing of fractional seconds, for example "15:59:59.989847074".

The function `as.POSIXct()` coerces objects into POSIXct *date-time* objects, with a numeric value representing the *moment of time* in seconds.


```
> library(HighFreq)
> # Read TAQ trade data from csv file
> taq <- data.table::fread(file="/Users/jerzy/Develop/lecture_slides/taq.csv")
> # Inspect the TAQ data in data.table format
> taq
> class(taq)
> colnames(taq)
> sapply(taq, class)
> symboln <- taq$SYM_ROOT[1]
> # Create date-time index
> datev <- paste(taq$DATE, taq$TIME_M)
> # Coerce date-time index to POSIXlt
> datev <- strptime(datev, "%Y%m%d %H:%M:%OS")
> class(datev)
> # Display more significant digits
> # options("digits")
> options(digits=20, digits.secs=10)
> last(datev)
> unclass(last(datev))
> as.numeric(last(datev))
> # Coerce date-time index to POSIXct
> datev <- as.POSIXct(datev)
> class(datev)
> last(datev)
> unclass(last(datev))
> as.numeric(last(datev))
> # Calculate the number of seconds
> as.numeric(last(datev)) - as.numeric(first(datev))
> # Calculate the number of ticks per second
> NROW(taq)/(6.5*3600)
> # Select TAQ data columns
> taq <- taq[, .(price=PRICE, volume=SIZE)]
```

Trading Volumes in High Frequency Data

The trading volumes represent the number of shares traded at a given price.

The histogram of the trading volumes shows that the highest frequencies of trades are for *round lots* - trades that are multiples of 100 shares.

There are also significant frequencies for *odd lots*, with small volumes of less than 100 shares.



```
> # Coerce trade ticks to xts series
> xlk <- xts::xts(taq[, .(price, volume)], datev)
> colnames(xlk) <- c("price", "volume")
> save(xlk, file="/Users/jerzy/Develop/data/xlk_tick_trades_2020031")
> # Plot histogram of the trading volumes
> hist(xlk$volume, main="Histogram of XLK Trading Volumes",
+      breaks=1e5, xlim=c(1, 400), xlab="number of shares")
```

Microstructure Noise in High Frequency Data


High frequency data contains *microstructure noise* in the form of *price spikes* and the *bid-ask bounce*.

Price spikes are single ticks with prices far away from the average.

Price spikes are often caused by data collection errors, but sometimes they represent actual trades with very large lot (trade) sizes.

The *bid-ask bounce* is the bouncing of traded prices between the bid and ask prices.

The *bid-ask bounce* creates an illusion of rapidly changing prices, while in reality the mid price is unchanged.



```
> # Plot dygraph
> dygraphs::dygraph(xlk$price, main="XLK Intraday Prices for 2020-03-16")
+ dyOptions(colors="blue", strokeWidth=1)
> # Plot in x11 window
> x11(width=6, height=5)
> quantmod::chart_Series(x=xlk$price, name="XLK Intraday Prices for 2020-03-16")
```

The Bid-ask Bounce of High Frequency Prices

The *bid-ask bounce* is the bouncing of traded prices between the bid and ask prices.


The *bid-ask bounce* is prominent at very high frequency time scales or in periods of low volatility.

The *bid-ask bounce* creates an illusion of rapidly changing prices, while in fact the mid price is constant.

The *bid-ask bounce* inflates the estimates of realized volatility, above the actual volatility.

The *bid-ask bounce* creates the appearance of mean reversion (negative autocorrelation), that isn't tradeable for most traders.

```
> pricev <- read.zoo(file="/Users/jerzy/Develop/lecture_slides/data,
+   header=TRUE, sep=",")
> pricev <- as.xts(pricev)
> dygraphs::dygraph(pricev$Close,
+   main="S&P500 Futures Prices Bid-Ask Bounce") %>%
+   dyOptions(colors="blue", strokeWidth=2)
```



Price Spikes And Trading Volumes in High Frequency Data

The number of the *price spikes* depends on the level of trading volumes, with the number decreasing with higher trading volumes.

The number of price spikes is lower for trade prices with larger trading volumes.



```
> # Plot dygraph of trade prices of at least 100 shares
> dygraphs::dygraph(xlk$price[xlk$volume >= 100, ],
+   main="XLK Prices for Trades of At Least 100 Shares") %>%
+   dyOptions(colors="blue", strokeWidth=1)
```


Removing Odd Lot Trades From TAQ Data

The trading volumes represent the number of shares traded at a given price.

The histogram of the trading volumes shows that the highest frequencies of trades are for *round lots* - trades that are multiples of 100 shares.

There are also significant frequencies for *odd lots*, with small volumes of less than 100 shares.

The *odd lot* ticks are often removed to reduce the size of the TAQ data.

Selecting only the large lot trades reduces microstructure noise (price spikes, bid-ask bounce) in high frequency data.

XLK Prices for Trades of At Least 100 Shares



```
> # Select the large trade lots of at least 100 shares
> dim(taq)
> tickb <- taq[taq$volume >= 100]
> dim(tickb)
> # Number of large lot ticks per second
> NROW(tickb)/(6.5*3600)
> # Plot histogram of the trading volumes
> hist(tickb$volume, main="Histogram of XLK Trading Volumes",
+      breaks=100000, xlim=c(1, 400), xlab="number of shares")
> # Save trade ticks with large lots
> data.table::fwrite(tickb, file="/Users/jerzy/Develop/data/xlk_tick_trades_2020-03-16_biglots.csv")
> # Coerce trade prices to xts
> xlbk <- xts::xts(tickb[, .(price, volume)], tickb$index)
> colnames(xlbk) <- c("price", "volume")
```

```
> # Plot dygraph of the large lots
> dygraphs::dygraph(xlbk$price,
+   main="XLK Prices for Trades of At Least 100 Shares") %>%
+   dyOptions(colors="blue", strokeWidth=1)
> # Plot the large lots
> x11(width=6, height=5)
> quantmod::chart_Series(x=xlk$price,
+   name="XLK Trade Ticks for 2020-03-16 (large lots only)")
```

Scrubbing Isolated Spikes In Stock Prices

Isolated price spikes can be identified using a *three-point filter* (tri-filter).

The *z-scores* are equal to the prices minus the mean of the neighboring prices, divided by their differences:

$$z_i = \frac{p_i - 0.5(p_{i-1} + p_{i+1}))}{p_{i-1} - p_{i+1}}$$

If the absolute value of the *z-score* exceeds the *threshold value* then it's classified as *bad data*, and it can be removed or replaced with the previous price.

The *tri-filter* can only remove isolated price spikes, but not consecutive price spikes.

XLK Intraday Prices for 2020-03-16



```
> # Load and plot intraday stock prices
> load("/Users/jerzy/Develop/lecture_slides/data/xlk_tick_trades_2
> pricev <- xlk$price
> dygraphs::dygraph(pricev, main="XLK Intraday Prices for 2020-03-
+   dyOptions(colors="blue", strokeWidth=1)
> # Calculate the lagged and advanced prices
> pricelag <- rutils::lagit(pricev)
> pricelag[1] <- pricelag[2]
> pricadv <- rutils::lagit(pricev, lag=-1)
> pricadv[NROW(pricadv)] <- pricadv[NROW(pricadv)-1]
> # Calculate the z-scores
> diff1 <- ifelse(abs(pricelag-pricadv) < 0.01, 0.01, abs(pricelag
> zscores <- (pricev - 0.5*(pricelag+pricadv))/diff1

> # Z-scores have very fat tails
> range(zscores); mad(zscores)
> madz <- mad(zscores[abs(zscores) > 0])
> hist(zscores, breaks=5000, xlim=c(-2*madz, 2*madz))
> # Scrub the price spikes
> threshv <- 5*madz ## Discrimination threshold
> indeks <- which(abs(zscores) > threshv)
> pricev[indeks] <- as.numeric(pricev[indeks-1])
> # Plot dygraph of the scrubbed prices
> dygraphs::dygraph(pricev, main="Scrubbed XLK Intraday Prices") %>
+   dyOptions(colors="blue", strokeWidth=1)
```

The Hampel Filter For Filtering Price Spikes

Price spikes in high frequency data can be identified using a *Hampel filter*.

The z-scores are equal to the prices minus the median of the prices, divided by the median absolute deviation (*MAD*) of prices:

$$z_i = \frac{p_i - \text{median}(p)}{\text{MAD}}$$

If the absolute value of the z-score exceeds the *threshold value* then it's classified as *bad data*, and it can be removed or replaced.

```
> # Calculate the centered Hampel filter to remove bad prices
> lookb <- 71 ## Look-back interval
> halfb <- lookb %/% 2 ## Half-back interval
> pricev <- xlk$price
> # Calculate the trailing median and MAD
> medianv <- HighFreq::roll_mean(pricev, lookb=lookb, method="nonp
> colnames(medianv) <- c("median")
> madv <- HighFreq::roll_var(pricev, lookb=lookb, method="nonparam
> # madv <- TTR::runMAD(pricev, n=lookb)
> # Center the median and the MAD
> medianv <- rutils::lagit(medianv, lagg=(-halfb), pad_zeros=FALSE)
> madv <- rutils::lagit(madv, lagg=(-halfb), pad_zeros=FALSE)
> # Calculate the Z-scores
> zscores <- ifelse(madv > 0, (pricev - medianv)/madv, 0)
> # Z-scores have very fat tails
> range(zscores); mad(zscores)
> madz <- mad(zscores[abs(zscores) > 0])
> hist(zscores, breaks=5000, xlim=c(-2*madz, 2*madz))
```

Scrubbed XLK Intraday Prices for 2020-03-16



```
> # Define discrimination threshold value
> threshv <- 6*madz
> # Identify good prices with small z-scores
> isgood <- (abs(zscores) < threshv)
> # Calculate the number of bad prices
> sum(!isgood)
> # Overwrite bad prices and calculate time series of scrubbed price
> pricec <- pricev
> pricec[!isgood] <- NA
> pricec <- zoo::na.locf(pricec)
> # Plot dygraph of the scrubbed prices
> dygraphs::dygraph(pricec, main="Scrubbed XLK Intraday Prices") %>
+   dyOptions(colors="blue", strokeWidth=1)
> # Plot using chart_Series()
> x11(width=6, height=5)
> quantmod::chart_Series(x=pricec,
+   name="Clean XLK Intraday Prices for 2020-03-16")
```

Classifying Data Outliers Using the Hampel Filter

The Hampel filter is a *classifier* which classifies the prices as either good or bad data points.

In order to measure the performance of the Hampel filter, we add price spikes to the clean prices, to see how accurately they're classified.

Let the *null hypothesis* be that the given price is a good data point.

A positive result corresponds to rejecting the *null hypothesis*, while a negative result corresponds to accepting the *null hypothesis*.

The classifications are subject to two different types of errors: *type I* and *type II* errors.

A *type I* error is the incorrect rejection of a TRUE *null hypothesis* (i.e. a "false positive"), when good data is classified as bad.

A *type II* error is the incorrect acceptance of a FALSE *null hypothesis* (i.e. a "false negative"), when bad data is classified as good.

```
> # Add 200 random price spikes to the clean prices
> set.seed(1121, "Mersenne-Twister", sample.kind="Rejection")
> nspikes <- 200
> nrows <- NROW(pricex)
> ispike <- logical(nrows)
> ispike[sample(x=nrows, size=nspikes)] <- TRUE
> priceb <- pricex
> priceb[ispike] <- priceb[ispike]*
+   sample(c(0.999, 1.001), size=nspikes, replace=TRUE)
> # Plot the bad prices and their medians
> medianv <- HighFreq::roll_mean(priceb, lookb=lookb, method="nonparametric")
> pricem <- cbind(priceb, medianv)
> colnames(pricem) <- c("prices with spikes", "median")
> dygraphs::dygraph(pricem, main="XLK Prices With Spikes") %>%
+   dyOptions(colors=c("red", "blue"))
> # Calculate the z-scores
> madv <- HighFreq::roll_var(priceb, lookb=lookb, method="nonparametric")
> zscores <- ifelse(madv > 0, (priceb - medianv)/madv, 0)
> # Z-scores have very fat tails
> range(zscores); mad(zscores)
> madz <- mad(zscores[abs(zscores) > 0])
> hist(zscores, breaks=10000, xlim=c(-4*madz, 4*madz))
> # Identify good prices with small z-scores
> threshv <- 3*madz
> isgood <- (abs(zscores) < threshv)
> # Calculate the number of bad prices
> sum(!isgood)
```

Confusion Matrix of a Binary Classification Model

A *binary classification model* categorizes cases based on its forecasts whether the *null hypothesis* is TRUE or FALSE.

The confusion matrix summarizes the performance of a classification model on a set of test data for which the actual values of the *null hypothesis* are known.

		Forecast	
		Null is FALSE	Null is TRUE
Actual	Null is FALSE	True Positive (sensitivity)	False Negative (type II error)
	Null is TRUE	False Positive (type I error)	True Negative (specificity)

```
> # Calculate the confusion matrix
> table(actual=!ispike, forecast=isgood)
> sum(!isgood)
> # FALSE positive (type I error)
> sum(!ispike & !isgood)
> # FALSE negative (type II error)
> sum(ispike & isgood)
```

Let the *null hypothesis* be that the given price is a good data point.

The *true positive rate* (known as the *sensitivity*) is the fraction of FALSE *null hypothesis* cases that are correctly classified as FALSE.

The *false negative rate* is the fraction of FALSE *null hypothesis* cases that are incorrectly classified as TRUE (*type II error*).

The sum of the *true positive* plus the *false negative* rate is equal to 1.

The *true negative rate* (known as the *specificity*) is the fraction of TRUE *null hypothesis* cases that are correctly classified as TRUE.

The *false positive rate* is the fraction of TRUE *null hypothesis* cases that are incorrectly classified as FALSE (*type I error*).

The sum of the *true negative* plus the *false positive* rate is equal to 1.


Receiver Operating Characteristic (ROC) Curve

The *ROC curve* is the plot of the *true positive rate*, as a function of the *false positive rate*, and illustrates the performance of a binary classifier.

The area under the *ROC curve* (AUC) measures the classification ability of a binary classifier.

```
> # Confusion matrix as function of threshold
> confun <- function(actualv, zscores, threshv) {
+   confmat <- table(actualv, (abs(zscores) < threshv))
+   confmat <- confmat / rowSums(confmat)
+   c(typeI=confmat[2, 1], typeII=confmat[1, 2])
+ } ## end confun
> confun(!ispike, zscores, threshv=threshv)
> # Define vector of discrimination thresholds
> threshv <- madz*seq(from=0.1, to=3.0, by=0.05)/2
> # Calculate the error rates
> errorr <- sapply(threshv, confun, actualv=!ispike, zscores=zscores)
> errorr <- t(errorr)
> rownames(errorr) <- threshv
> errorr <- rbind(c(1, 0), errorr)
> errorr <- rbind(errorr, c(0, 1))
> # Calculate the area under the ROC curve (AUC)
> truepos <- (1 - errorr[, "typeII"])
> truepos <- (truepos + rutils::lagit(truepos))/2
> falsepos <- rutils::diffit(errorr[, "typeI"])
> abs(sum(truepos*falsepos))
```

ROC Curve for Hampel Classifier



```
> # Plot ROC curve for Hampel classifier
> plot(x=errorr[, "typeI"], y=1-errorr[, "typeII"],
+      xlab="FALSE positive rate", ylab="TRUE positive rate",
+      xlim=c(0, 1), ylim=c(0, 1),
+      main="ROC Curve for Hampel Classifier",
+      type="l", lwd=3, col="blue")
> abline(a=0.0, b=1.0, lwd=3, col="orange")
```


Filtering Bad Data From Daily Stock Prices

Daily stock prices can also contain bad data points consisting of mostly single, isolated spikes in prices.

The number of false positives may be too high, so the Hampel filter parameters (the look-back interval and the threshold) need adjustment.

For example, the VXX has only one bad price (on 2010-11-08), but the Hampel filter identifies many more than that (which are false positives).

```
> # Load log VXX prices
> load("/Users/jerzy/Develop/lecture_slides/data/pricevxx.RData")
> nrows <- NROW(pricev)
> # Calculate the centered Hampel filter for VXX
> lookb <- 7 ## Look-back interval
> halfb <- lookb %/% 2 ## Half-back interval
> medianv <- HighFreq::roll_mean(pricev, lookb=lookb, method="nonp
> medianv <- rutils::lagit(medianv, lagg=(-halfb), pad_zeros=FALSE)
> madv <- HighFreq::roll_var(pricev, lookb=lookb, method="nonparam
> madv <- rutils::lagit(madv, lagg=(-halfb), pad_zeros=FALSE)
> zscores <- ifelse(madv > 0, (pricev - medianv)/madv, 0)
> range(zscores); mad(zscores)
> madz <- mad(zscores[abs(zscores) > 0])
> hist(zscores, breaks=100, xlim=c(-3*madz, 3*madz))
> # Define discrimination threshold value
> threshv <- 9*madz
> # Calculate the good prices
> isgood <- (abs(zscores) < threshv)
> sum(!isgood)
> # Dates of the bad prices
> zoo::index(pricev[!isgood])
```



```
> # Calculate the false positives
> falsep <- !isgood
> falsep[which(zoo::index(pricev) == as.Date("2010-11-08"))] <- FALSE
> # Plot dygraph of the prices with bad prices
> datam <- cbind(pricev, zscores)
> colnames(datam)[2] <- "ZScores"
> colv <- colnames(datam)
> dygraphs::dygraph(datam, main="VXX Prices With Z-Scores and False
+ dyAxis("y", label=colv[1], independentTicks=TRUE) %>%
+ dyAxis("y2", label=colv[2], independentTicks=TRUE) %>%
+ dySeries(name=colv[1], axis="y", strokeWidth=1, col="blue") %>%
+ dySeries(name=colv[2], axis="y2", strokeWidth=1, col="red") %>%
+ dyEvent(zoo::index(pricev[falsep]), label=rep("false", sum(falsep)),
+ dyEvent(zoo::index(pricev["2010-11-08"]), label="true", strokeP
```

draft: Filtering Combined Spikes From Stock Prices

The narrow Hampel filter isn't very good anyway

The narrow Hampel filter using the median of 3 prices can only identify single isolated spikes.

But sometimes several bad prices occur in a row, one after another.

The narrow Hampel filter cannot identify multiple bad prices in a row, and will therefore produce false negatives (bad prices identified as good).

```
> # Add single isolated spike to the prices
> priceb <- pricev
> priceb["2017-11-20"] <- 1.2*priceb["2017-11-20"]
> # Calculate the Z-scores
> medianv <- HighFreq::roll_mean(priceb, lookb=lookb, method="nonparametric", lagg=(-halfb), pad_zeros=FALSE)
> medianv <- rutils::lagit(medianv, lagg=(-halfb), pad_zeros=FALSE)
> madv <- HighFreq::roll_var(priceb, lookb=lookb, method="nonparametric", lagg=(-halfb), pad_zeros=FALSE)
> madv <- rutils::lagit(madv, lagg=(-halfb), pad_zeros=FALSE)
> zscores <- ifelse(madv > 0, (priceb - medianv)/madv, 0)
> madz <- mad(zscores[abs(zscores) > 0])
> # Calculate the number of bad prices
> threshv <- 9*madz
> isgood <- (abs(zscores) < threshv)
> sum(!isgood)
> zoo::index(priceb[!isgood])
> # Add two neighboring spikes to the prices
> priceb <- pricev
> priceb["2017-11-20"] <- 1.2*priceb["2017-11-21"]
> priceb["2017-11-21"] <- 1.2*priceb["2017-11-21"]
> # Calculate the Z-scores
> medianv <- HighFreq::roll_mean(priceb, lookb=lookb, method="nonparametric", lagg=(-halfb), pad_zeros=FALSE)
> medianv <- rutils::lagit(medianv, lagg=(-halfb), pad_zeros=FALSE)
> madv <- HighFreq::roll_var(priceb, lookb=lookb, method="nonparametric", lagg=(-halfb), pad_zeros=FALSE)
> madv <- rutils::lagit(madv, lagg=(-halfb), pad_zeros=FALSE)
> zscores <- ifelse(madv > 0, (priceb - medianv)/madv, 0)
> madz <- mad(zscores[abs(zscores) > 0])
> # Calculate the number of bad prices
> isgood <- (abs(zscores) < threshv)
> sum(!isgood)
> zoo::index(priceb[!isgood])
```


Scrubbing Bad Stock Prices

Bad stock prices can be scrubbed (replaced) with the previous good price.

But it's incorrect to replace bad prices with the average of the previous good price and the next good price, since that would cause data snooping.

```
> # Replace bad stock prices with the previous good prices
> priceg <- pricev
> priceg[!isgood] <- NA
> priceg <- zoo::na.locf(priceg)
> # Calculate the Z-scores
> medianv <- HighFreq::roll_mean(priceg, lookb=lookb, method="nonpara
> medianv <- rutils::lagit(medianv, lagg=(-halfb), pad_zeros=FALSE)
> madv <- HighFreq::roll_var(priceg, lookb=lookb, method="nonparamet
> madv <- rutils::lagit(madv, lagg=(-halfb), pad_zeros=FALSE)
> zscores <- ifelse(madv > 0, (priceg - medianv)/madv, 0)
> madz <- mad(zscores[abs(zscores) > 0])
> # Calculate the number of bad prices
> threshv <- 9*madz
> isgood <- (abs(zscores) < threshv)
> sum(!isgood)
> zoo::index(priceg[!isgood])
```

Scrubbed VXX Prices With False Positives



```
> # Calculate the false positives
> falsep <- !isgood
> falsep[which(zoo::index(pricev) == as.Date("2010-11-08"))] <- FALSE
> # Plot dygraph of the prices with bad prices
> dygraphs::dygraph(priceg, main="Scrubbed VXX Prices With False Pos
+   dyEvent(zoo::index(priceg[falsep]), label=rep("false", sum(falsep)
+   dyOptions(colors="blue", strokeWidth=1)
```

ROC Curve for Daily Hampel Classifier


In order to measure the performance of the Hampel filter, we add price spikes to the clean prices, to see how accurately they're classified.

The performance of the Hampel noise classification model depends on the length of the look-back time interval.

The optimal *look-back interval* and *threshold value* can be determined using *cross-validation*.

```
> # Add 200 random price spikes to the clean prices
> set.seed(1121, "Mersenne-Twister", sample.kind="Rejection")
> nspikes <- 200
> ispike <- logical(nrows)
> ispike[sample(x=nrows, size=nspikes)] <- TRUE
> priceb <- priceg
> priceb[ispike] <- priceb[ispike]*
+   sample(c(0.99, 1.01), size=nspikes, replace=TRUE)
> # Calculate the Z-scores
> medianv <- HighFreq::roll_mean(priceb, lookb=lookb, method="nonparametric")
> medianv <- rutils::lagit(medianv, lagg=(-halfb), pad_zeros=FALSE)
> madv <- HighFreq::roll_var(priceb, lookb=lookb, method="nonparam")
> madv <- rutils::lagit(madv, lagg=(-halfb), pad_zeros=FALSE)
> zscores <- ifelse(madv > 0, (priceb - medianv)/madv, 0)
> madz <- mad(zscores[abs(zscores) > 0])
> # Define vector of discrimination thresholds
> threshv <- madz*seq(from=0.1, to=3.0, by=0.05)/2
> # Calculate the error rates
> errorrr <- sapply(threshv, confun, actualv=!ispike, zscores=zscores)
> errorrr <- t(errorrr)
> rownames(errorrr) <- threshv
> errorrr <- rbind(c(1, 0), errorrr)
> errorrr <- rbind(errorrr, c(0, 1))
```

ROC Curve for Daily Hampel Classifier



```
> # Calculate the area under the ROC curve (AUC)
> truepos <- (1 - errorrr[, "typeII"])
> truepos <- (truepos + rutils::lagit(truepos))/2
> falsepos <- rutils::diffit(errorrr[, "typeI"])
> abs(sum(truepos*falsepos))
> # Plot ROC curve for Hampel classifier
> plot(x=errorrr[, "typeI"], y=1-errorrr[, "typeII"],
+      xlab="FALSE positive rate", ylab="TRUE positive rate",
+      xlim=c(0, 1), ylim=c(0, 1),
+      main="ROC Curve for Daily Hampel Classifier",
+      type="l", lwd=3, col="blue")
> abline(a=0.0, b=1.0, lwd=3, col="orange")
```

Aggregating TAQ Data to OHLC

The *data table* columns can be *aggregated* over categories (factors) defined by one or more columns passed to the "by" argument.

Multiple *data table* columns can be referenced by passing a list of names specified by the dot `.` operator.

The function `round.POSIXt()` rounds date-time objects to seconds, minutes, hours, days, months or years.

The function `as.POSIXct()` coerces objects to class `POSIXct`.

```
> # Round time index to seconds
> tickg[, zoo::index := as.POSIXct(round.POSIXt(index, "secs"))]
> # Aggregate to OHLC by seconds
> ohlc <- tickg[, .(open=first(price), high=max(price), low=min(price), close=last(price)), by=index)
> # Round time index to minutes
> tickg[, zoo::index := as.POSIXct(round.POSIXt(index, "mins"))]
> # Aggregate to OHLC by minutes
> ohlc <- tickg[, .(open=first(price), high=max(price), low=min(price), close=last(price)), by=index)
```



```
> # Coerce OHLC prices to xts
> ohlc <- xts::xts(ohlc[, -"index"], ohlc$index)
> # Plot dygraph of the OHLC prices
> dygraphs::dygraph(ohlc[, -5], main="XLK Trade Ticks for 2020-03-16")
+   dyCandlestick()
> # Plot the OHLC prices
> x11(width=6, height=5)
> quantmod::chart_Series(x=ohlc, TA="add_Vo()",
+   name="XLK Trade Ticks for 2020-03-16 (OHLC)")
```

draft: Trade and Quote (TAQ) Data

Trade and Quote (TAQ) Data

The NYSE TAQ set is 'consolidated' meaning that the trades and quotes come from many exchanges, including the NASDAQ system. TAQ has all trades and quotes from the Consolidated Tape / Consolidated Quote, which includes data from the regional exchanges. (See

<https://www.tradearca.com/marketdata/intermarket.asp> for a brief discussion of CTA/CQ and OTC/UTP.)

<https://wrds->

www.wharton.upenn.edu/pages/support/research-wrds/research-guides/wrds-taq-faqs/

High frequency data is typically formatted as either Trade and Quote (TAQ) data, or *Open-High-Low-Close* (OHLC) data.

Trade and Quote (TAQ) data contains intraday trades and quotes on exchange-traded stocks and futures.

The TAQ data is spaced irregularly in time, with data recorded each time a new trade or quote arrives.

Each row of TAQ data contains both the quote and trade prices, and the corresponding quote size or trade volume.

Each row of TAQ data contains both the quote and trade prices, and the corresponding quote size or trade volume: *Bid.Price*, *Bid.Size*, *Ask.Price*, *Ask.Size*, *Trade.Price*, *Volume*.

```
> # Load package HighFreq
> library(HighFreq)
> # Or load the high frequency data file directly:
> # symbolv <- load("/Users/jerzy/Develop/R/HighFreq/data/hf_data.R")
> head(HighFreq::SPY_TAQ)
```

	Bid.Price	Bid.Size	Ask.Price	Ask.Size	Trade.Price
2014-05-02 00:00:01	188	1	189	15	
2014-05-02 08:00:01	188	1	189	5	
2014-05-02 08:00:02	189	1	189	5	
2014-05-02 08:01:13	188	1	189	5	
2014-05-02 08:01:29	188	1	189	5	
2014-05-02 08:01:52	189	2	189	5	

```
> head(HighFreq::SPY)
```

	SPY.Open	SPY.High	SPY.Low	SPY.Close	SPY.Volume
2008-01-02 09:31:00	147	147	147	147	591203
2008-01-02 09:32:00	147	147	147	147	385457
2008-01-02 09:33:00	147	147	147	147	343700
2008-01-02 09:34:00	147	147	147	147	863418
2008-01-02 09:35:00	147	147	147	147	457500
2008-01-02 09:36:00	147	147	147	147	416708

```
> tail(HighFreq::SPY)
```

	SPY.Open	SPY.High	SPY.Low	SPY.Close	SPY.Volume
2014-05-19 15:56:00	189	189	189	189	321321
2014-05-19 15:57:00	189	189	189	189	299474
2014-05-19 15:58:00	189	189	189	189	261357
2014-05-19 15:59:00	189	189	189	189	435886
2014-05-19 16:00:00	189	189	189	189	1824185
2014-05-19 16:01:00	189	189	189	189	63920

Open-High-Low-Close (OHLC) Data

Open-High-Low-Close (OHLC) data contains intraday trade prices and trade volumes.

OHLC data is evenly spaced in time, with each row containing the *Open*, *High*, *Low*, *Close* prices, and the trade *Volume*, recorded over the past time interval (called a *bar* of data).

The *Open* and *Close* prices are the first and last trade prices recorded in the time bar.

The *High* and *Low* prices are the highest and lowest trade prices recorded in the time bar.

The *Volume* is the total trading volume recorded in the time bar.

The *OHLC* data format provides a way of efficiently compressing *TAQ* data, while preserving information about price levels, volatility (range), and trading volumes.

In addition, evenly spaced *OHLC* data allows for easier analysis of multiple time series, since the prices for different assets are given at the same moments in time.

```
> # Load package HighFreq
> library(HighFreq)
> head(HighFreq::SPY)
```

	SPY.Open	SPY.High	SPY.Low	SPY.Close	SPY.Volume
2008-01-02 09:31:00	147	147	147	147	591203
2008-01-02 09:32:00	147	147	147	147	385457
2008-01-02 09:33:00	147	147	147	147	343700
2008-01-02 09:34:00	147	147	147	147	863418
2008-01-02 09:35:00	147	147	147	147	457500
2008-01-02 09:36:00	147	147	147	147	416708

Plotting High Frequency OHLC Data

Aggregating high frequency *TAQ* data into *OHLC* format with lower periodicity allows for data compression while maintaining some information about volatility.

```
> # Load package HighFreq
> library(HighFreq)
> # Define symbol
> symboln <- "SPY"
> # Load OHLC data
> dirout <- "/Users/jerzy/Develop/data/minutes/"
> symboln <- load(file.path(dirout, paste0(symboln, ".RData")))
> interval <- "2013-11-11 09:30:00/2013-11-11 10:30:00"
> chart_Series(SPY[interval], name=symboln)
```

The package *HighFreq* contains both *TAQ* data and *Open-High-Low-Close (OHLC)* data.

If you are not able to install package *HighFreq* then download the file *hf_data.RData* from the NYU share drive and load it.



Package *HighFreq* for Managing High Frequency Data

The package *HighFreq* contains functions for managing high frequency time series data, such as:

- converting *TAQ* data to *OHLC* format,
- chaining and joining time series,
- scrubbing bad data,
- managing time zones and aligning time indices,
- aggregating data to lower frequency (periodicity),
- calculating rolling aggregations (VWAP, Hurst exponent, etc.),
- calculating seasonality aggregations,
- estimating volatility, skewness, and higher moments,

```
> # Install package HighFreq from github
> devtools::install_github(repo="algoquant/HighFreq")
> # Load package HighFreq
> library(HighFreq)
> # Get documentation for package HighFreq
> # Get short description
> packageDescription(HighFreq)
> # Load help page
> help(package=HighFreq)
> # List all datasets in HighFreq
> data(package=HighFreq)
> # List all objects in HighFreq
> ls("package:HighFreq")
> # Remove HighFreq from search path
> detach("package:HighFreq")
```

draft: Package HighFreq for Managing High Frequency Data

Package HighFreq contains functions for managing high frequency *TAQ* and *OHLC* market data:

- reading and writing data from files,
- managing time zones and aligning indices,
- chaining and joining time series,
- scrubbing bad data points,
- converting *TAQ* data to *OHLC* format,
- aggregating data to lower frequency,

HighFreq is inspired by the package *highfrequency*, and follows many of its conventions.

HighFreq depends on packages *xts*, *quantmod*, *lubridate*, and *caTools*.

The function `scrub_agg()` scrubs a single day of *TAQ* data, aggregates it, and converts it to *OHLC* format.

The function `save_scrub_agg()` loads, scrubs, aggregates, and binds multiple days of *TAQ* data for a single symbol, and saves the *OHLC* time series to a single **.RData* file.

```
> # Install package HighFreq from github
> install.packages("devtools")
> library(devtools)
> install_github(repo="algoquant/HighFreq")
> # Load package HighFreq
> library(HighFreq)
> # Set data directories
> dirin <- "/Users/jerzy/Develop/data/hfreq/src/"
> dirout <- "/Users/jerzy/Develop/data/hfreq/scrub/"
> # Define symbol
> symboln <- "SPY"
> # Load a single day of TAQ data
> symboln <- load(file.path(dirin, paste0(symboln, "/2014.05.02."), s
> # Scrub, aggregate single day of TAQ data to OHLC
> ohlc_data <- scrub_agg(taq_data=get(symboln))
> # Aggregate TAQ data for the symbol, and save to file
> HighFreq::save_scrub_agg(symboln,
+   dirin=dirin,
+   dirout=dirout,
+   period="minutes")
```


Datasets in Package *HighFreq*

The package *HighFreq* contains several high frequency time series, in *xts* format, stored in a file called `hf_data.RData`:

- a time series called `SPY.TAQ`, containing a single day of *TAQ* data for the *SPY* ETF.
- three time series called `SPY`, `TLT`, and `VXX`, containing intraday 1-minute *OHLC* price bars for the *SPY*, *TLT*, and *VXX* ETFs.

Even after the *HighFreq* package is loaded, its datasets aren't loaded into the workspace, so they aren't listed in the workspace.

That's because the datasets in package *HighFreq* are set up for *lazy loading*, which means they can be called as if they were loaded, even though they're not loaded into the workspace.

The datasets in package *HighFreq* can be loaded into the workspace using the function `data()`.

The data is set up for *lazy loading*, so it doesn't require calling `data(hf_data)` to load it into the workspace before calling it.

```
> # Load package HighFreq
> library(HighFreq)
> # You can see SPY when listing objects in HighFreq
> ls("package:HighFreq")
> # You can see SPY when listing datasets in HighFreq
> data(package=HighFreq)
> # But the SPY dataset isn't listed in the workspace
> ls()
> # HighFreq datasets are lazy loaded and available when needed
> head(HighFreq::SPY)
> # Load all the datasets in package HighFreq
> data(hf_data)
> # HighFreq datasets are now loaded and in the workspace
> head(HighFreq::SPY)
```

Distribution of High Frequency Returns

High frequency returns exhibit *large negative skewness* and *very large kurtosis* (leptokurtosis), or fat tails.


Student's *t-distribution* has fat tails, so it fits high frequency returns much better than the normal distribution.

The function `fitdistr()` from package *MASS* fits a univariate distribution into a sample of data, by performing *maximum likelihood* optimization.

The function `hist()` calculates and plots a histogram, and returns its data *invisibly*.

```
> # Calculate SPY percentage returns
> ohlc <- HighFreq::SPY
> nrow <- NROW(ohlc)
> closep <- log(quantmod::Cl(ohlc))
> retp <- rutils::diffit(closep)
> colnames(retp) <- "SPY"
> # Standardize raw returns to make later comparisons
> retp <- (retp - mean(retp))/sd(retp)
> # Calculate moments and perform normality test
> sapply(c(var=2, skew=3, kurt=4), function(x) sum(retp^x)/nrow)
> tseries::jarque.bera.test(retp)
> # Fit SPY returns using MASS::fitdistr()
> optiml <- MASS::fitdistr(retp, densfun="t", df=2)
> loc <- optiml$estimate[1]
> scalev <- optiml$estimate[2]
```

Distribution of High Frequency SPY Returns



```
> # Plot histogram of SPY returns
> histp <- hist(retp, col="lightgrey", mgp=c(2, 1, 0),
+ xlab="returns (standardized)", ylab="frequency", xlim=c(-3, 3),
+ breaks=1e3, freq=FALSE, main="Distribution of High Frequency SPY")
> # lines(density(retp, bw=0.2), lwd=3, col="blue")
> # Plot t-distribution function
> curve(expr=dt((x-loc)/scalev, df=2)/scalev,
+ type="l", lwd=3, col="red", add=TRUE)
> # Plot the Normal probability distribution
> curve(expr=dnorm(x, mean=mean(retp),
+ sd=sd(retp)), add=TRUE, lwd=3, col="blue")
> # Add legend
> legend("topright", inset=0.05, bty="n",
+ leg=c("t-distr", "normal"), y.intersp=0.5,
+ lwd=6, lty=1, col=c("red", "blue"))
```


Distribution of Aggregated High Frequency Returns

The distribution of returns depends on the sampling frequency.

High frequency returns aggregated to a lower periodicity become less negatively skewed and less fat tailed, and closer to the normal distribution.

The function `xts::to.period()` converts a time series to a lower periodicity (for example from hourly to daily periodicity).

```
> # Hourly SPY percentage returns
> closep <- log(Cl(xts::to.period(x=ohlcv, period="hours")))
> retsh <- rutils::diffit(closep)
> retsh <- (retsh - mean(retsh))/sd(retsh)
> # Daily SPY percentage returns
> closep <- log(Cl(xts::to.period(x=ohlcv, period="days")))
> retld <- rutils::diffit(closep)
> retld <- (retld - mean(retld))/sd(retld)
> # Calculate moments
> sapply(list(minutely=retp, hourly=retsh, daily=retld),
+ function(rets) {sapply(c(var=2, skew=3, kurt=4),
+ function(x) mean(rets^x))
+ }) ## end sapply
```




```
> # Plot densities of SPY returns
> plot(density(retp, bw=0.4), xlim=c(-3, 3),
+      lwd=3, mgp=c(2, 1, 0), col="blue",
+      xlab="returns (standardized)", ylab="frequency",
+      main="Density of High Frequency SPY Returns")
> lines(density(retsh, bw=0.4), lwd=3, col="green")
> lines(density(retld, bw=0.4), lwd=3, col="red")
> # Add legend
> legend("topright", inset=0.05, bty="n",
+       leg=c("minutely", "hourly", "daily"), y.intersp=0.5,
+       lwd=6, lty=1, col=c("blue", "green", "red"))
```

Estimating Rolling Volatility of High Frequency Returns

The volatility of high frequency returns can be inflated by large overnight returns.

The large overnight returns can be scaled down by dividing them by the overnight time interval.

```
> # Calculate rolling volatility of SPY returns
> ret2013 <- retp["2013-11-11/2013-11-15"]
> # Calculate rolling volatility
> lookb <- 11 ## Look-back interval
> endd <- seq_along(ret2013)
> startp <- c(rep_len(1, lookb),
+   endd[1:(NROW(endd)-lookb)])
> endd[endd < lookb] <- lookb
> vol_rolling <- sapply(seq_along(endd),
+   function(it) sd(ret2013[startp[it]:endd[it]]))
> vol_rolling <- xts::xts(vol_rolling, zoo::index(ret2013))
> # Extract time intervals of SPY returns
> indeks <- c(60, diff(xts::index(ret2013)))
> head(indeks)
> table(indeks)
> # Scale SPY returns by time intervals
> ret2013 <- 60*ret2013/indeks
> # Calculate scaled rolling volatility
> vol_scaled <- sapply(seq_along(endd),
+   function(it) sd(ret2013[startp[it]:endd[it]]))
> vol_rolling <- cbind(vol_rolling, vol_scaled)
> vol_rolling <- na.omit(vol_rolling)
> sum(is.na(vol_rolling))
> sapply(vol_rolling, range)
```



```
> # Plot rolling volatility
> x11(width=6, height=5)
> themev <- chart_theme()
> themev$col$line.col <- c("blue", "red")
> chart_Series(vol_rolling, theme=themev,
+   name="Rolling Volatility with Overnight Spikes")
> legend("topright", legend=colnames(vol_rolling),
+   inset=0.1, bg="white", lty=1, lwd=6, y.intersp=0.5,
+   col=themev$col$line.col, bty="n")
```


Daily Volume and Volatility of High Frequency Returns

Trading volumes typically rise together with market price volatility.

The function `apply.daily()` from package `xts` applies functions to time series over daily periods.

The function `calc_var_ohlc()` from package `HighFreq` calculates the variance of an *OHLC* time series using range estimators.

```
> # Volatility of SPY
> sqrt(HighFreq::calcvar_ohlc(ohlc))
> # Daily SPY volatility and volume
> volatd <- sqrt(xts::apply.daily(ohlc, FUN=calcvar_ohlc))
> colnames(volatd) <- ("SPY_volatility")
> volumv <- quantmod::Vo(ohlc)
> volumd <- xts::apply.daily(volumv, FUN=sum)
> colnames(volumd) <- ("SPY_volume")
> # Plot SPY volatility and volume
> datav <- cbind(volatd, volumd)["2008/2009"]
> colv <- colnames(datav)
> dygraphs::dygraph(datav,
+   main="SPY Daily Volatility and Trading Volume") %>%
+   dyAxis("y", label=colv[1], independentTicks=TRUE) %>%
+   dyAxis("y2", label=colv[2], independentTicks=TRUE) %>%
+   dySeries(name=colv[1], axis="y", col="red", strokeWidth=3) %>%
+   dySeries(name=colv[2], axis="y2", col="blue", strokeWidth=3)
```




Beta of Volume vs Volatility of High Frequency Returns

As a general empirical rule, the *trading volume* v in a given time period is roughly proportional to the *volatility* of the returns σ : $v \propto \sigma$.

The regression of the *log trading volume* versus the *log volatility* fails the *Durbin-Watson test* for the autocorrelation of residuals.

But the regression of the *differences* passes the *Durbin-Watson test*.

```
> # Regress log of daily volume vs volatility
> datav <- log(cbind(volumd, volatd))
> colv <- colnames(datav)
> dframe <- as.data.frame(datav)
> formulav <- as.formula(paste(colv, collapse="~"))
> regmod <- lm(formulav, data=dframe)
> # Durbin-Watson test for autocorrelation of residuals
> lmtest::dwtest(regmod)
> # Regress diff log of daily volume vs volatility
> dframe <- as.data.frame(rutils::diffit(datav))
> regmod <- lm(formulav, data=dframe)
> lmtest::dwtest(regmod)
> summary(regmod)
> plot(formulav, data=dframe, main="SPY Daily Trading Volume vs Volatility (log scale)")
> abline(regmod, lwd=3, col="red")
> mtext(paste("beta =", round(coef(regmod)[2], 3)), cex=1.2, lwd=3, side=2, las=2, adj=(-0.5), padj=(-7))
```




Hourly Trading Volume vs Volatility

Hourly aggregations of high frequency data also support the rule that the *trading volume* is roughly proportional to the *volatility* of the returns: $v \propto \sigma$.

```
> # 60 minutes of data in lookb interval
> lookb <- 60 ## Look-back interval
> vol2013 <- volumv["2013"]
> ret2013 <- retp["2013"]
> # Define end points with beginning stub
> nrow <- NROW(ret2013)
> nagg <- nrow %/% lookb
> endd <- nrow-lookb*nagg + (0:nagg)*lookb
> startp <- c(1, endd[1:(NROW(endd)-1)])
> # Calculate SPY volatility and volume
> datav <- sapply(seq_along(endd), function(it) {
+   endp <- startp[it]:endd[it]
+   c(volume=sum(vol2013[endp]),
+     volatility=sd(ret2013[endp]))
+ }) ## end sapply
> datav <- t(datav)
> datav <- rutils::diffit(log(datav))
> dframe <- as.data.frame(datav)
```

SPY Hourly Trading Volume vs Volatility (log scale)



```
> formulav <- as.formula(paste(colnames(datav), collapse=""))
> regmod <- lm(formulav, data=dframe)
> lmtest::dwtest(regmod)
> summary(regmod)
> plot(formulav, data=dframe,
+   main="SPY Hourly Trading Volume vs Volatility (log scale)")
> abline(regmod, lwd=3, col="red")
> mtext(paste("beta =", round(coef(regmod)[2], 3)), cex=1.2, lwd=3,
```

High Frequency Returns in Trading Time

The *trading time* (volume clock) is the time measured by the level of *trading volume*, with the *volume clock* running faster in periods of higher *trading volume*.


The time-dependent volatility of high frequency returns (*heteroskedasticity*) produces their *leptokurtosis* (large kurtosis, or fat tails).

The returns can be divided by the *square root of the trading volumes* to obtain scaled returns over equal trading volumes.

But the returns should not be divided by very small volumes below a certain threshold.

The scaled returns have a smaller *skewness* and *kurtosis*, and they also have even higher autocorrelations than unscaled returns.

Density of Volume-scaled High Frequency SPY Returns



```
> # Scale returns using volume (volume clock)
> retsc <- ifelse(volumv > 1e4, retp/sqrt(volumv), 0)
> retsc <- retsc/sd(retsc)
> # Calculate moments of scaled returns
> nrow <- NROW(retp)
> sapply(list(retp=retp, retsc=retsc),
+   function(rets) {sapply(c(skew=3, kurt=4),
+     function(x) sum((rets/sd(rets))^x)/nrow)
+ }) ## end sapply
```

```
> x11(width=6, height=5)
> par(mar=c(3, 3, 2, 1), oma=c(1, 1, 1, 1))
> # Plot densities of SPY returns
> plot(density(retp), xlim=c(-3, 3),
+   lwd=3, mgp=c(2, 1, 0), col="blue",
+   xlab="returns (standardized)", ylab="frequency",
+   main="Density of Volume-scaled High Frequency SPY Returns")
> lines(density(retsc, bw=0.4), lwd=3, col="red")
> curve(expr=dnorm, add=TRUE, lwd=3, col="green")
> # Add legend
> legend("topright", inset=0.05, bty="n", y.intersp=0.5,
+   leg=c("minutely", "scaled", "normal"),
+   lwd=6, lty=1, col=c("blue", "red", "green"))
```


Autocorrelations of High Frequency Returns

The *Ljung-Box* test, tests if the autocorrelations of a time series are *statistically significant*.

The *null hypothesis* of the *Ljung-Box* test is that the autocorrelations are equal to zero.

The *Ljung-Box* statistic is small for time series that have *statistically insignificant* autocorrelations.

The function `Box.test()` calculates the *Ljung-Box* test and returns the test statistic and its p-value.

For *minutely SPY* returns, the *Ljung-Box* statistic is large and its *p*-value is very small, so we can conclude that *minutely SPY* returns have statistically significant autocorrelations.

For *scaled minutely SPY* returns, the *Ljung-Box* statistic is even larger, so its autocorrelations are even more statistically significant.

SPY returns aggregated to longer time intervals are less autocorrelated.

```
> # Ljung-Box test for minutely SPY returns
> Box.test(retp, lag=10, type="Ljung")
> # Ljung-Box test for daily SPY returns
> Box.test(retd, lag=10, type="Ljung")
> # Ljung-Box test statistics for scaled SPY returns
> sapply(list(retp=retp, retsc=retsc),
+   function(rets) {
+     Box.test(rets, lag=10, type="Ljung")$statistic
+ }) ## end sapply
> # Ljung-Box test statistics for aggregated SPY returns
> sapply(list(minutely=retp, hourly=retsh, daily=retd),
+   function(rets) {
+     Box.test(rets, lag=10, type="Ljung")$statistic
+ }) ## end sapply
```

The level of the autocorrelations depends on the sampling frequency, with higher frequency returns having more significant negative autocorrelations.

As the returns are aggregated to a lower periodicity, they become less autocorrelated, with daily returns having almost insignificant autocorrelations.


Partial Autocorrelations of High Frequency Returns

High frequency minutely *SPY* returns have statistically significant negative autocorrelations.


SPY returns *scaled* by the trading volumes have even more significant negative autocorrelations.

```
> # Set plot parameters
> x11(width=6, height=8)
> par(mar=c(4, 4, 2, 1), oma=c(0, 0, 0, 0))
> layout(matrix(c(1, 2), ncol=1), widths=c(6, 6), heights=c(4, 4))
> # Plot the partial autocorrelations of minutely SPY returns
> pacf1 <- pacf(as.numeric(retp), lag=10,
+   xlab="lag", ylab="partial autocorrelation", main="")
> title("Partial Autocorrelations of Minutely SPY Returns", line=1)
> # Plot the partial autocorrelations of scaled SPY returns
> pacfs <- pacf(as.numeric(retsc), lag=10,
+   xlab="lag", ylab="partial autocorrelation", main="")
> title("Partial Autocorrelations of Scaled SPY Returns", line=1)
> # Calculate the sums of partial autocorrelations
> sum(pacf1$acf)
> sum(pacfs$acf)
```

Partial Autocorrelations of Minutely SPY Return



Partial Autocorrelations of Scaled SPY Return



Market Liquidity, Trading Volume and Volatility

Market illiquidity is defined as the market price impact resulting from supply-demand imbalance.

Market liquidity $\mathcal{L}$ is proportional to the square root of the trading volume v divided by the price volatility σ :


$$\mathcal{L} \sim \frac{\sqrt{v}}{\sigma}$$

Market illiquidity spiked during the May 6, 2010 *flash crash*.

Research suggests that market crashes are caused by declining market liquidity:

Donier et al., Why Do Markets Crash?


```
> # Calculate market illiquidity
> liquidityv <- sqrt(volumd)/volatd
> # Plot market illiquidity
> x11(width=6, height=7) ; par(mfrow=c(2, 1))
> themev <- chart_theme()
> themev$col$line.col <- c("blue")
> chart_Series(liquidityv["2010"], theme=themev,
+   name="SPY Liquidity in 2010", plot=FALSE)
> themev$col$line.col <- c("red")
> chart_Series(volatd["2010"],
+   theme=themev, name="SPY Volatility in 2010")
```



Intraday Seasonality of Volume and Volatility

The volatility and trading volumes are typically higher at the beginning and end of the trading sessions.

```
> # Calculate intraday time index with hours and minutes
> datev <- format(zoo::index(retp), "%H:%M")
> # Aggregate the mean volume
> volumagg <- tapply(X=volumv, INDEX=datev, FUN=mean)
> volumagg <- drop(volumagg)
> # Aggregate the mean volatility
> volagg <- tapply(X=retp^2, INDEX=datev, FUN=mean)
> volagg <- sqrt(drop(volagg))
> # Coerce to xts
> datev <- as.POSIXct(paste(Sys.Date(), names(volumagg)))
> volumagg <- xts::xts(volumagg, datev)
> volagg <- xts::xts(volagg, datev)
> # Plot seasonality of volume and volatility
> x11(width=6, height=7) ; par(mfrow=c(2, 1))
> themev <- chart_theme()
> themev$col$line.col <- c("blue")
> chart_Series(volumagg[c(-1, -NROW(volumagg))], theme=themev,
+   name="Intraday Seasonality of SPY Volume", plot=FALSE)
> themev$col$line.col <- c("red")
> chart_Series(volagg[c(-1, -NROW(volagg))], theme=themev,
+   name="Intraday Seasonality of SPY Volatility")
```




Intraday Seasonality of Liquidity and Volatility

Market liquidity is typically the highest at the end of the trading session, and the lowest at the beginning.

The end of day spike in trading volumes and liquidity is driven by computer-driven investors liquidating their positions.


```
> # Calculate market liquidity
> liquidv <- sqrt(volumagg)/volagg
> # Plot intraday seasonality of market liquidity
> x11(width=6, height=7) ; par(mfrow=c(2, 1))
> themev <- chart_theme()
> themev$col$line.col <- c("blue")
> chart_Series(liquidv[c(-1, -NROW(liquidv))], theme=themev,
+   name="Intraday Seasonality of SPY Liquidity", plot=FALSE)
> themev$col$line.col <- c("red")
> chart_Series(volagg[c(-1, -NROW(volagg))], theme=themev,
+   name="Intraday Seasonality of SPY Volatility")
```



draft: Daily Volatility and Hurst Exponent

The Hurst exponent typically moves higher with higher market price volatility, and is above 0.5 with high volatility.


```
> chart_Series(roll_sum(volatd, 10)[-1:10])/10, name=paste(symbol,
> chart_Series(roll_sum(hurstd, 10)[-1:10])/10, name=paste(symbol,
> abline(h=0.5, col="blue", lwd=2)
```



draft: Intraday Seasonality of Hurst Exponent and Volatility

The Hurst exponent typically moves higher with higher market price volatility, and is above 0.5 with high volatility.

```
> # Intraday seasonality of Hurst exponent
> interval <- "2013"
> season_hurst <- season_ality(hurst_ohlc(ohlc=SPY[interval, 1:4]))
> season_hurst <- season_hurst[-(nrow(season_hurst))]
> colnames(season_hurst) <- paste0(colname(get(symboln)), ".season_")
> themev <- chart_theme()
> themev$format.labels <- "%H:%M"
> chobj <- chart_Series(x=season_hurst,
+   name=paste(colnames(season_hurst),
+   "intraday seasonality"), theme=themev,
+   plot=FALSE)
> ylim <- chobj$getYlim()
> ylim[[2]] <- structure(c(ylim[[2]][1],
+   ylim[[2]][2]), fixed=TRUE)
> chobj$set_ylim(ylim)
> plot(chobj)
> abline(h=0.5, col="blue", lwd=2)
> # Intraday seasonality of volatility
> season_var <- season_ality(vol_ohlc(ohlc=SPY))
```




draft: Estimating Skew From OHLC Data

The function `skew_ohlc()` from package `HighFreq` calculates a skew-like indicator:

$$s^3 = \frac{1}{n} \sum_{i=1}^n ((H_i - O_i)(H_i - C_i)(H_i - 0.5(O_i + C_i)) + (L_i - O_i)(L_i - C_i)(L_i - 0.5(O_i + C_i)))$$

The function `roll_agg_ohlc()` aggregates rolling, volume weighted moment estimators.

```
> library(rutils) ## Load package rutils
> # Rolling variance
> varv <- roll_agg_ohlc(ohlc=SPY, agg_fun="vol_ohlc")
> # Rolling skew
> skewn <- roll_agg_ohlc(ohlc=SPY, agg_fun="skew_ohlc")
> skewn <- skewn/(varv)^(1.5)
> skewn[1, ] <- 0
> skewn <- zoo::na.locf(skewn)
> interval <- "2013-11-11/2013-11-15"
> chart_Series(varv[interval],
+             name=paste(symboln, "variance"))
> chart_Series(skewn[interval],
+             name=paste(symboln, "Skew"),
+             ylim=c(-1, 1))
```



draft: Daily Volatility and Skew From OHLC Data

The function `agg_ohlc()` calculates a statistical estimator over an *OHLC* time series.


```
> # Daily variance and skew
> volatd <- xts::apply.daily(x=HighFreq::SPY, FUN=agg_ohlc,
+   agg_fun="vol_ohlc")
> colnames(volatd) <- paste0(symboln, ".var")
> daily_skew <- xts::apply.daily(x=HighFreq::SPY, FUN=agg_ohlc,
+   agg_fun="skew_ohlc")
> daily_skew <- daily_skew/(volatd)^(1.5)
> colnames(daily_skew) <- paste0(symboln, ".skew")
> interval <- "2013-06-01/"
> chart_Series(volatd[interval],
+   name=paste(symboln, "variance"))
> chart_Series(daily_skew[interval],
+   name=paste(symboln, "skew"))
```



draft: Regression of Skews Versus Returns

A regression of lagged skews versus returns appears to be statistically significant, especially in periods of high volatility during the financial crisis of 2008–09.

```
> retp <- calc_rets(xts_data=SPY)
> skewn <- skew_ohlc(log_ohlc=log(SPY[, -5]))
> colnames(skewn) <- paste0(symboln, ".skew")
> skewl <- rutils::lag_it(skewn)
> skewl[1, ] <- 0
> datav <- cbind(retp[, 1], sign(skewl))
> formulav <- as.formula(paste(colnames(datav)[1],
+   paste(paste(colnames(datav)[-1],
+     collapse=" + "), "~ 1"), sep="~"))
> formulav
> regmod <- lm(formulav, data=datav)
> summary(regmod)$coef
> summary(lm(formulav, data=datav["/2011-01-01/"]))$coef
> summary(lm(formulav, data=datav["2011-01-01/"]))$coef
```



```
> interval <- "2013-12-01/"
> plot(formulav, data=datav[interval],
+   xlim=c(-2e-09, 2e-09),
+   cex=0.6, xlab="skew", ylab="rets")
> abline(regmod, col="blue", lwd=2)
```

draft: Contrarian Strategy Using Skew Oscillator

The contrarian skew trading strategy involves long or short positions in a single unit of stock, that is opposite to the sign of the skew.

Skew is calculated over one-minute bars, and trades are executed in the following period.

The contrarian strategy shows good hypothetical performance before transaction costs, and since it's a liquidity providing strategy, should have very low transaction costs.

The contrarian strategy is hyperactive, trading almost 46% of the time in each period.

```
> # Lag the skew to get positions
> posit <- -sign(skew1)
> posit[1, ] <- 0
> # Cumulative PnL
> pnl <- cumsum(posit*retp[, 1])
> # Calculate frequency of trades
> 50*sum(abs(sign(skewn)-sign(skew1)))/nrow(skewn)
> # Calculate transaction costs
> bidask <- 0.001 ## 10 bps for liquid ETFs
> bidask*sum(abs(sign(skewn)-sign(skew1)))
```



```
> chart_Series(pnl[endpoints(pnl, on="hours"), ],
+ name=paste(symboln, "Contrarian Skew Strategy PnL"))
```


draft: Volume-Weighted Average Price Indicator

The Volume-Weighted Average Price (VWAP) is an indicator used for trend following strategies.

The fast-moving VWAP is calculated over a short look-back interval, while the slow-moving VWAP is calculated over a longer interval.

The trend following reverses direction when the fast-moving VWAP crosses the slow-moving one.


```
> vwapf <- vwapv(xtsv=SPY, lookb=70)
> vwaps <- vwapv(xtsv=SPY, lookb=225)
> vwapd <- vwapf - vwaps
> colnames(vwapd) <- paste0(symboln, ".vwap")
> interval <- "2010-05-05/2010-05-07"
> invisible(chart_Series(x=C1(SPY[interval]), name=paste(symboln, "plus
> invisible(add_TA(vwapf[interval], on=1, col="red", lwd=2))
> invisible(add_TA(vwaps[interval], on=1, col="blue", lwd=2))
> invisible(add_TA(vwapd[interval] > 0, on=-1, col="lightgreen", border
> add_TA(vwapd[interval] < 0, on=-1, col="lightgrey", border="lightgrey"
```



draft: Regression of VWAP Versus Returns

A regression of the *VWAP* indicator versus returns appears to be statistically significant.

```
> lag_vwap <- rutils::lag_it(vwapd)
> lag_vwap[1, ] <- 0
> datav <- cbind(retp[, 1], sign(lag_vwap))
> formulav <- as.formula(paste(colnames(datav)[1],
+   paste(paste(colnames(datav)[-1],
+     collapse=" + "), "- 1"), sep="--"))
> formulav
> regmod <- lm(formulav, data=datav)
> summary(regmod)$coef
> summary(lm(formulav, data=datav["/2011-01-01/"]))$coef
> summary(lm(formulav, data=datav["2011-01-01/"]))$coef
```



```
> interval <- "2013-12-01/"
> plot(formulav, data=cbind(retp[, 1], lag_vwap)[interval],
+   cex=0.6, xlab="skew", ylab="rets")
> abline(regmod, col="blue", lwd=2)
```

draft: Trend Following Strategy Using VWAP


The trend following trading strategy involves long or short positions in a single unit of stock, that is equal to the sign of the VWAP indicator.

The VWAP indicator is calculated over one-minute bars, and trades are executed in the following period.

The trend following strategy shows good hypothetical performance before transaction costs.

The trend following strategy is infrequent, trading only 0.56% of the time in each period.

```
> # Cumulative PnL
> pnl <- cumsum(sign(lag_vwap)*ret[p, 1])
> # Calculate frequency of trades
> 50*sum(abs(sign(vwapd)-sign(lag_vwap)))/nrow(vwapd)
> # Calculate transaction costs
> bidask <- 0.001 ## 10 bps for liquid ETFs
> bidask*sum(abs(sign(vwapd)-sign(lag_vwap)))
```



```
> chart_Series(
+   pnl[endpoints(pnl, on="hours"), ],
+   name=paste(symboln, "VWAP Trend Following Strategy PnL"))
```

draft: Estimating Hurst Exponent From OHLC Data

The Hurst exponent is a measure of long-term memory of a time series, and is related to its autocorrelation:


$$\mathbb{E}\left[\frac{\max(p) - \min(p)}{\hat{\sigma}}\right] = t^H$$

$H = 0.5$ for Brownian motion (no autocorrelations),
 $H > 0.5$ for positive autocorrelations,
 $H < 0.5$ for negative autocorrelations.

The function `hurst_ohlcl()` from package `HighFreq` calculates a Hurst-like indicator:

$$H = \frac{1}{n} \sum_{i=1}^n \log\left(\frac{H_i - L_i}{\text{abs}(C_i - O_i)}\right)$$

The function `agg_ohlcl()` calculates a statistical estimator over an OHLC time series.



```
> # Daily Hurst exponents
> hurstd <- xts::apply.daily(x=HighFreq::SPY, FUN=agg_ohlcl, agg_fun=
> colnames(hurstd) <- paste(colname(get(symboln)), ".Hurst")
> chart_Series(roll_sum(hurstd, 10)[-1:10])/10, name=paste(symboln,
> abline(h=0.5, col="blue", lwd=2)
```

draft: Conclusion

Open questions:

- is there an interaction between volatility, volume and skew?
- does kurtosis also have predictive value for market direction?
- how can higher moments (skew, kurtosis) help predict market crashes?
- what is the persistence of market anomalies over time?
- what is relationship between returns and cross-section of skew?
- This presentation is available on GitHub: <https://github.com/algoquant/presentations>

Acknowledgements:

- *Snowfall Systems* provides the *PortfolioEffect* system for:
 - real-time high frequency market data aggregations and risk metrics,
 - real-time portfolio analytics and optimization,
 - portfolio hosting,
 - <https://www.portfolioeffect.com/>
- Brian Peterson for Thomson Reuters tick data,
- Brian Peterson, Joshua Ulrich, and Jeffrey Ryan for packages *xts*, *quantmod*, *PerformanceAnalytics*, and *TTR*,



Package *IBrokers* for Using Interactive Brokers

Interactive Brokers (IB) is a brokerage company which provides an API for computer-driven trading, and it also provides extensive documentation:

[Interactive Brokers Main Page](#)
[Interactive Brokers Documentation](#)
[Interactive Brokers Course](#)

Disclosure: I do have a personal account with Interactive Brokers.

But I do not have any other relationship with Interactive Brokers, and I do not endorse or recommend them.

Warning: Active trading is extremely risky, and most people lose money.

I advise not to trade with your own capital!

The package *IBrokers* contains R functions for executing IB commands using the Interactive Brokers API.

The package *IBrokers* has extensive [documentation](#).

```
> # Install package IBrokers
> install.packages("IBrokers")
> # Load package IBrokers
> library(IBrokers)
> # Get documentation for package IBrokers
> # Get short description
> packageDescription("IBrokers")
> # Load help page
> help(package="IBrokers")
> # List all datasets in "IBrokers"
> data(package="IBrokers")
> # List all objects in "IBrokers"
> ls("package:IBrokers")
> # Remove IBrokers from search path
> detach("package:IBrokers")
> # Install package IBrokers2
> devtools::install_github(repo="algoquant/IBrokers2")
```

The package *IBrokers2* is derived from package *IBrokers*, and contains additional functions for executing real time trading strategies via the Interactive Brokers API.

Configuring the IB Trader Workstation (TWS)

Connecting to Interactive Brokers via the API requires being logged into the IB Trader Workstation (TWS) or the IB Gateway (IBG).



You should [Download the TWS](#) and install it.

Read about the [TWS initial setup](#).

The TWS settings must be configured to enable the API: *File → Global Configuration → API*

The TWS Java heap size should be increased to 1.5 GB: *File → Global Configuration → General*

Read more about the [required TWS memory allocation](#) and about [how to change the TWS heap size](#).



Connecting to Interactive Brokers via the API

Connecting to Interactive Brokers via the API requires being logged into the IB Trader Workstation (*TWS*) or the IB Gateway (*IBG*).

Interactive Brokers provides extensive [API Documentation](#).

The functions `twConnect()` and `ibgConnect()` open a connection to the Interactive Brokers API, via either the *TWS* or the *IBG*.

The parameter `port` should be assigned to the value of the *socket port* displayed in *TWS* or *IBG*.

The function `twDisconnect()` closes the Interactive Brokers API connection.

```
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twConnect(port=7497)
> # Or connect to IB Gateway
> # connib <- ibgConnect(port=4002)
> # Check connection
> IBrokers::isConnected(connib)
> # Close the Interactive Brokers API connection
> IBrokers::twDisconnect(connib)
```

Downloading Account Information from Interactive Brokers

The function `reqAccountUpdates()` returns a list with the account information (requires an account code).

The first element of the list contains the account dollar balances, while the remaining elements contain contract information.

The function `twPortfolioValue()` returns a data frame with commonly used account fields, such as the contract names, net positions, and realized and unrealized profits and losses (pnl's).

```
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twConnect(port=7497)
> # Or connect to IB Gateway
> # connib <- ibgConnect(port=4002)
> # Download account information from IB
> ac_count <- "DU1215081"
> ib_account <- IBrokers::reqAccountUpdates(conn=connib,
+                                           acctCode=ac_count)
> # Extract account balances
> balance_s <- ib_account[[1]]
> balance_s$AvailableFunds
> # Extract contract names, net positions, and profits and losses
> IBrokers::twPortfolioValue(ib_account)
> # Close the Interactive Brokers API connection
> IBrokers::twDisconnect(connib)
```

Defining Contracts Using Package *IBrokers*

The function `twsequity()` defines a stock contract (IB contract object).

The functions `twsfuture()` and `twscurrency()` define futures and currency contracts.

The function `reqContractDetails()` returns a list with information on the IB instrument.

The package *twbinstrument* contains utility functions for enhancing the package *IBrokers*.

To define an IB contract, first look it up by keyword in the online [IB Contract and Symbol Database](#), and find the *Conid* for that instrument.

Enter the *Conid* into the function `twbinstrument::getContract()`, which will return the IB contract object for that instrument.

Interactive Brokers provides more information about financial contracts here: [IB Traded Products](#)

```
> # Define AAPL stock contract (object)
> contractobj <- IBrokers::twsequity("AAPL", primary="SMART")
> # Define CHF currency contract
> contractobj <- IBrokers::twscurrency("CHF", currency="USD")
> # Define S&P Emini future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="ES",
+   exch="GLOBEX", expiry="201906")
> # Define 10yr Treasury future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="ZN",
+   exch="ECBOT", expiry="201906")
> # Define euro currency future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="EUR",
+   exch="GLOBEX", expiry="201906")
> # Define Gold future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="GC",
+   exch="NYMEX", expiry="201906")
> # Define Oil future January 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="QM",
+   exch="NYMEX", expiry="201901")
> # Test if contract object is correct
> IBrokers::is.twsContract(contractobj)
> # Get list with instrument information
> IBrokers::reqContractDetails(conn=connib, Contract=contractobj)
> # Install the package twbinstrument
> install.packages("twbinstrument", repos="http://r-forge.r-project.org")
> # Define euro future using getContract() and Conid
> contractobj <- twbinstrument::getContract("317631411")
> # Get list with instrument information
> IBrokers::reqContractDetails(conn=connib, Contract=contractobj)
```

Defining *VIX* Futures Contracts

VIX futures have both monthly and weekly contracts, with the monthly contracts being the most liquid.

In order to uniquely specify a *VIX* futures contract, the parameter "local" should be passed into the function `twFuture()`, with the local security name.

For example, *VXV8* is the local security name (symboln) for the monthly *VIX* futures contract expiring on October 17th, 2018.

VX40V8 is the local security name (symboln) for the weekly *VIX* futures contract expiring on October 3rd, 2018.

The function `reqContractDetails()` returns a list with information on the IB instrument.

VIX futures are traded on the *CFE* (CBOE Futures Exchange): <http://cfe.cboe.com/>

```
> # Define VIX monthly and weekly futures June 2019 contract
> symboln <- "VIX"
> contractobj <- IBrokers::twFuture(symbol=symboln,
+   exch="CFE", expiry="201906")
> # Define VIX monthly futures June 2019 contract
> contractobj <- IBrokers::twFuture(symbol=symboln,
+   local="VXV8", exch="CFE", expiry="201906")
> # Define VIX weekly futures October 3rd 2018 contract
> contractobj <- IBrokers::twFuture(symbol=symboln,
+   local="VX40V8", exch="CFE", expiry="201906")
> # Get list with instrument information
> IBrokers::reqContractDetails(conn=connib,
+   Contract=contractobj)
```

Downloading Historical Daily Data from Interactive Brokers

The function `reqHistoricalData()` downloads historical data to a .csv file.

The historical *daily* bar data fields are "Open", "High", "Low", "Close", "Volume", "WAP", "Count". ("WAP" is the weighted average price.)

Interactive Brokers provides more information about historical market data:

[IB Historical Market Data](#)

[IB Historical Bar Data Fields](#)

```
> # Define S&P Emini futures June 2019 contract
> symboln <- "ES"
> contractobj <- IBrokers::twFuture(symbol=symboln,
+   exch="GLOBEX", expiry="201906")
> # Open file for data download
> dirn <- "/Users/jerzy/Develop/data/vix"
> dir.create(dirn)
> filen <- file.path(dirn, paste0(symboln, "201906.csv"))
> filec <- file(filen, open="w")
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twConnect(port=7497)
> # Write header to file
> cat(paste(paste(symboln, c("Index", "Open", "High", "Low", "Close",
+   Contract=contractobj,
+   barSize="1 day", duration="6 M",
+   file=filec)
> # Close data file
> close(filec)
> # Close the Interactive Brokers API connection
> IBrokers::twDisconnect(connib)
```

Downloading Historical Intraday Data for a Portfolio

Historical data for a portfolio of symbols can be downloaded in a loop.

The historical *intraday* bar data fields are "Open", "High", "Low", "Close", "Volume", "WAP", "XTRA", "Count". ("WAP" is the weighted average price.)


```
> # Define IB contract objects for stock symbols
> symbolv <- c("AAPL", "F", "MSFT")
> contractv <- lapply(symbolv, IBrokers::twseEquity, primary="SMART")
> names(contractv) <- symbolv
> # Open file connections for data download
> dirn <- "/Users/jerzy/Develop/data/vix"
> filev <- file.path(dirn, paste0(symbolv, format(Sys.time(), format
> fileconn <- lapply(filev, function(filen) file(filen, open="w"))
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twseConnect(port=7497)
> # Download historical 1-minute bar data to files
> for (it in 1:NROW(symbolv)) {
+   symboln <- symbolv[it]
+   filec <- fileconn[[it]]
+   contractobj <- contractv[[it]]
+   cat("Downloading data for: ", symboln, "\n")
+   ## Write header to file
+   cat(paste(paste(symboln, c("Index", "Open", "High", "Low", "Close",
+     IBrokers::reqHistoricalData(conn=connib,
+       Contract=contractobj,
+       barSize="1 min", duration="2 D",
+       file=filec)
+   Sys.sleep(10) ## 10s pause to avoid IB pacing violation
+ } ## end for
> # Close data files
> for (filec in fileconn) close(filec)
> # Close the Interactive Brokers API connection
> IBrokers::twseDisconnect(connib)
```


Downloading Historical Data for Expired Futures Contracts

Historical data for expired futures contracts can be downloaded using `reqHistoricalData()` with the appropriate expiry date, and the parameter `include_expired="1"`.

For example, *ESM8* is the symbol for the *S&P500* emini futures expiring in June 2018.

```
> # Define S&P Emini futures June 2018 contract
> symboln <- "ES"
> contractobj <- IBrokers::twFuture(symbol=symboln,
+   include_expired="1",
+   exch="GLOBEX", expiry="201806")
> # Open file connection for ESM8 data download
> filen <- file.path(dirn, paste0(symboln, "M8.csv"))
> filec <- file(filen, open="w")
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twConnect(port=7497)
> # Download historical data to file
> IBrokers::reqHistoricalData(conn=connib,
+   Contract=contractobj,
+   barSize="1 day", duration="2 Y",
+   file=filec)
> # Close data file
> close(filec)
> # Close the Interactive Brokers API connection
> IBrokers::twDisconnect(connib)
```



```
> # Load OHLC data and coerce it into xts series
> pricev <- data.table::fread(filen)
> data.table::setDF(pricev)
> pricev <- xts::xts(pricev[, 2:6],
+   order.by=as.Date(as.POSIXct.numeric(pricev[, 1],
+   tz="America/New_York", origin="1970-01-01")))
> colnames(pricev) <- c("Open", "High", "Low", "Close", "Volume")
> # Plot OHLC data in x11 window
> chart_Series(x=pricev, TA="add_Vo()",
+   name="S&P500 ESM8 futures")
> # Plot dygraph
> dygraphs::dygraph(pricev[, 1:4], main="S&P500 ESM8 futures") %>%
+   dyCandlestick()
```

draft: Downloading Continuous Contract Futures Data

Note: Continuous futures require passing the parameter: `contract.secType = "CONFUT"`, which isn't currently included in package *IBrokers*.

http://interactivebrokers.github.io/tws-api/basic_contracts.html

A continuous futures contract is a synthetic time series of prices, created by splicing together prices from several futures contracts with different expiration dates.

At any point in time, the price of the continuous futures is equal to the most liquid contract times a normalization factor.

When the next consecutive contract becomes more liquid, then the continuous futures price is rolled over to that contract.

The continuous price is multiplied by a normalization factor when the contract is rolled, to remove jumps caused by the shape of the futures curve.

So the continuous futures prices are not equal to the past futures prices.

Futures contracts trade at different prices (because of the futures convenience yield).

This cause price jumps between the currently expiring futures contract and the next futures contract. A continuous futures contract adjusts the prices to remove these jumps and time differences to create an artificial price series.

```
> # Define S&P Emini futures June 2018 contract
> symboln <- "ES"
> contractobj <- IBrokers::twsFuture(symbol=symboln,
+   include_expired="1",
+   exch="GLOBEX", expiry="201806")
> # Open file connection for data download
> dirn <- "/Users/jerzy/Develop/data/vix"
> dir.create(dirn)
> filen <- file.path(dirn, paste0(symboln, ".csv"))
> filec <- file(filen, open="w")
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twsConnect(port=7497)
> # Download historical data to file
> IBrokers::reqHistoricalData(conn=connib,
+   Contract=contractobj,
+   barSize="1 day", duration="6 M",
+   file=filec)
> # Close data file
> close(filec)
> # Close the Interactive Brokers API connection
> IBrokers::twsDisconnect(connib)
```

Downloading Live *TAQ* Data from Interactive Brokers

The function `reqMktData()` downloads live (real-time) trades and quotes (*TAQ*) data from Interactive Brokers.

The function `eWrapper()` formats the real-time market events (trades and quotes), so they can be displayed or saved to a file.

The method `eWrapper.MktData.CSV()` formats the real-time *TAQ* data so it can be saved to a `.csv` file.

The real-time *TAQ* data fields are *BidSize*, *BidPrice*, *AskPrice*, *AskSize*, *Last*, *LastSize*, *Volume*.

BidPrice is the quoted bid price, *AskPrice* is the quoted offer price, and *Last* is the most recent traded price.

The *TAQ* data is spaced irregularly in time, with data recorded each time a new trade or quote arrives.

```
> # Define S&P Emini futures June 2019 contract
> symboln <- "ES"
> contractobj <- IBrokers::twFuture(symbol=symboln,
+   exch="GLOBEX", expiry="201906")
> # Open file connection for data download
> dirn <- "/Users/jerzy/Develop/data/vix"
> # Dir.create(dirn)
> filen <- file.path(dirn, paste0(symboln, "_taq_live.csv"))
> filec <- file(filen, open="w")
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twConnect(port=7497)
> # Download live data to file
> IBrokers::reqMktData(conn=connib,
+   Contract=contractobj,
+   eventWrapper=eWrapper.MktData.CSV(1),
+   file=filec)
> # Close data file
> close(filec)
> # Close the Interactive Brokers API connection
> IBrokers::twDisconnect(connib)
```

Downloading Live *OHLC* Data from Interactive Brokers

The function `reqRealTimeBars()` downloads live (real-time) *OHLC* market data from Interactive Brokers.

The method `eWrapper.RealTimeBars.CSV()` formats the real-time *OHLC* data so it can be saved to a `.csv` file.

Interactive Brokers by default only provides 5-second bars of real-time prices (but it also provides historical data at other frequencies).

```
> # Define S&P Emini futures June 2019 contract
> symboln <- "ES"
> contractobj <- IBrokers::twFuture(symbol=symboln,
+   exch="GLOBEX", expiry="201906")
> # Open file connection for data download
> dirn <- "/Users/jerzy/Develop/data/vix"
> # Dir.create(dirn)
> filen <- file.path(dirn, paste0(symboln, "_ohlc_live.csv"))
> filec <- file(filen, open="w")
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twConnect(port=7497)
> # Download live data to file
> IBrokers::reqRealTimeBars(conn=connib,
+   Contract=contractobj, barSize="1",
+   eventWrapper=eWrapper.RealTimeBars.CSV(1),
+   file=filec)
> # Close the Interactive Brokers API connection
> IBrokers::twDisconnect(connib)
> # Close data file
> close(filec)
> # Load OHLC data and coerce it into xts series
> library(data.table)
> pricev <- data.table::fread(filen)
> pricev <- xts::xts(pricev[, paste0("V", 2:6)],
+   as.POSIXct.numeric(as.numeric(pricev[, V1]), tz="America/New_York"))
> colnames(pricev) <- c("Open", "High", "Low", "Close", "Volume")
> # Plot OHLC data in x11 window
> x11()
> chart_Series(x=pricev, TA="add_Vo()",
+   name="S&P500 ESM9 futures")
> # Plot dygraph
> library(dygraphs)
> dygraphs::dygraph(pricev[, 1:4], main="S&P500 ESM9 futures") %>%
+   dyCandlestick()
```

Downloading Live *OHLC* Data For Multiple Contracts

Live *OHLC* data can be downloaded for multiple instruments simultaneously.

This requires passing a *list* of contracts and file connections to `reqRealTimeBars()`, and also passing the number of contracts to `eWrapper.RealTimeBars.CSV()`.

The bar prices for each contract are written into a separate file.

```
> library(IBrokers)
> # Define list of S&P futures and 10yr Treasury contracts
> contractv <- list(ES=IBrokers::twFuture(symbol="ES", exch="GLOBEX"),
+                 ZN=IBrokers::twFuture(symbol="ZN", exch="ECBOT", exchLoc="U.S."))
> # Open the file connection for storing the bar data
> dirn <- "/Users/jerzy/Develop/data/vix"
> filev <- file.path(dirn, paste0(c("ES", "ZN"), format(Sys.time(), "%Y%m%d")))
> fileconn <- lapply(filev, function(file) file(file, open="w"))
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twConnect(port=7497)
> # Download live data to file
> IBrokers::reqRealTimeBars(conn=connib,
+                           Contract=contractv,
+                           barSize="1", useRTH=FALSE,
+                           eventWrapper=eWrapper.RealTimeBars.CSV(NROW=contracts,
+                           file=fileconn))
> # Close the Interactive Brokers API connection
> IBrokers::twDisconnect(connib)
> # Close data files
> for (filec in fileconn)
+   close(filec)
> library(data.table)
> # Load ES futures June 2019 contract and coerce it into xts series
> pricev <- data.table::fread(filev[1])
> pricev <- xts::xts(pricev[, paste0("V", 2:6)],
+                   as.POSIXct.numeric(as.numeric(pricev[, V1]), tz="America/New_York"))
> colnames(pricev) <- c("Open", "High", "Low", "Close", "Volume")
> # Plot dygraph
> library(dygraphs)
> dygraphs::dygraph(pricev[, 1:4], main="S&P500 ESM9 futures") %>%
+   dyCandlestick()
> # Load ZN 10yr Treasury futures June 2019 contract
> pricev <- data.table::fread(filev[2])
> pricev <- xts::xts(pricev[, paste0("V", 2:6)],
+                   as.POSIXct.numeric(as.numeric(pricev[, V1]), tz="America/New_York"))
> colnames(pricev) <- c("Open", "High", "Low", "Close", "Volume")
> # Plot dygraph
> dygraphs::dygraph(pricev[, 1:4], main="7Y 10yr Treasury futures") %>%
+   dyCandlestick()
```

Event Processing Using `eWrapper()` and `twcCALLBACK()`

Functions which process real-time market events (like `reqMktData()` and `reqRealTimeBars()`) rely on the functions `eWrapper()` and `twcCALLBACK()`.


The function `eWrapper()` creates an *eWrapper* object, consisting of a data environment and handler (methods) for formatting and adding new data to it.

The function `reqMktData()` accepts an *eWrapper* object and passes it into `twcCALLBACK()`.

The function `twcCALLBACK()` processes market events by calling functions `readBin()` and `processMsg()` in a `while()` loop.

The TWS broadcasts real-time market events in the form of character strings, which are captured by `readBin()`.

The character strings are then parsed by `processMsg()`, and copied to the data environment of the *eWrapper* object.



Placing Market Trade Orders on *TWS*

The function `reqIds()` requests a trade order ID from Interactive Brokers *TWS*.

The function `twOrder()` creates a trade order object.

The parameter `orderType` specifies the type of trade order: market order (MKT), limit order (LMT), etc.

Each trade order requires its own ID generated by `reqIds()`.

The function `placeOrder()` places a trade order.

```
> # Define Japanese yen currency futures June 2019 contract 6JZ8
> contractobj <- IBrokers::twFuture(symbol="JPY", exch="GLOBEX", ex
> # Connect to Interactive Brokers TWS
> connib <- IBrokers::twConnect(port=7497)
> IBrokers::reqContractDetails(conn=connib, Contract=contractobj)
> # Request trade order ID
> order_id <- IBrokers::reqIds(connib)
> # Create buy market order object
> ib_order <- IBrokers::twOrder(order_id,
+   orderType="MKT", action="BUY", totalQuantity=1)
> # Place trade order
> IBrokers::placeOrder(connib, contractobj, ib_order)
> # Execute sell market order
> order_id <- IBrokers::reqIds(connib)
> ib_order <- IBrokers::twOrder(order_id,
+   orderType="MKT", action="SELL", totalQuantity=1)
> IBrokers::placeOrder(connib, contractobj, ib_order)
> # Execute buy market order
> order_id <- IBrokers::reqIds(connib)
> ib_order <- IBrokers::twOrder(order_id,
+   orderType="MKT", action="BUY", totalQuantity=1)
> IBrokers::placeOrder(connib, contractobj, ib_order)
```

Placing and Cancelling Limit Trade Orders on *TWS*

The function `tw$Order()` with `orderType=LMT` creates a limit trade order object.

The function `cancelOrder()` with the parameter `orderId=order_id` cancels a trade order with the ID equal to `order_id`.

```
> # Request trade order ID
> order_id <- IBrokers::reqIds(connib)
> # Create buy limit order object
> ib_order <- IBrokers::tw$Order(order_id, orderType="LMT",
+   lmtPrice="1.1511", action="BUY", totalQuantity=1)
> # Place trade order
> IBrokers::placeOrder(connib, contractobj, ib_order)
> # Cancel trade order
> IBrokers::cancelOrder(connib, order_id)
> # Execute sell limit order
> order_id <- IBrokers::reqIds(connib)
> ib_order <- IBrokers::tw$Order(order_id, orderType="LMT",
+   lmtPrice="1.1512", action="SELL", totalQuantity=1)
> IBrokers::placeOrder(connib, contractobj, ib_order)
> # Cancel trade order
> IBrokers::cancelOrder(connib, order_id)
> # Close the Interactive Brokers API connection
> IBrokers::tw$Disconnect(connib)
```


Placing Limit Orders in a Programmatic Loop

The function `reqRealTimeBars()` creates data objects and then calls `tw$CALLBACK()` and passes the objects to `tw$CALLBACK()`.

`tw$CALLBACK()` calls `processMsg()` in a loop, and passes the `eWrapper()` to it.

`processMsg()` performs if-else statements and calls `eWrapper` methods.

`eWrapper()` creates an environment for data, and methods (handlers) for formatting the data.

`eWrapper.RealTimeBars.CSV()` calls `eWrapper()` to create a closure (function) and returns it.

Closures allow creating mutable states for persistent data objects.

Limit trade orders can be placed inside a modified `eWrapper()` function.

```
> eWrapper_realtimebars <- function(n = 1) {
+   eW <- eWrapper_new(NULL)
+   ## eW <- IBrokers::eWrapper(NULL)
+   eW$assign.Data("data", rep(list(structure(.xts(matrix(rep(NA, rep(1, n)), nrow = 100),
+     eW$realtimeBars <- function(curMsg, msg, timestamp, file, ...) {
+       id <- as.numeric(msg[2])
+       file <- file[[id]]
+       data <- eW$get.Data("data")
+       attr(data[[id]], "index") <- as.numeric(msg[3])
+       nr.data <- NROW(data[[id]])
+       ## Write to file
+       cat(paste(msg[3], msg[4], msg[5], msg[6], msg[7], msg[8], msg[9], msg[10], "\n"),
+         file=file, append=TRUE)
+       ## Write to console
+       ## eW$countn <- eW$countn + 1
+       eW$assign.Data("countn", eW$get.Data("countn") + 1)
+       cat(paste0("countn=", eW$get.Data("countn"), "\tOpen=", msg[4],
+         ## cat(paste0("Open=", msg[4], "\tHigh=", msg[5], "\tLow=", msg[6], "\tClose=", msg[7], "\n"),
+         ## Trade
+         ## Cancel previous trade orders
+       buy_id <- eW$get.Data("buy_id")
+       sell_id <- eW$get.Data("sell_id")
+       if (buy_id>0) IBrokers::cancelOrder(connib, buy_id)
+       if (sell_id>0) IBrokers::cancelOrder(connib, sell_id)
+       ## Execute buy limit order
+       buy_id <- IBrokers::reqIds(connib)
+       buy_order <- IBrokers::tw$Order(buy_id, orderType="LMT",
+         lmtPrice=msg[6]-0.25, action="BUY",
+       IBrokers::placeOrder(connib, contractobj, buy_order)
+       ## Execute sell limit order
+       sell_id <- IBrokers::reqIds(connib)
+       sell_order <- IBrokers::tw$Order(sell_id, orderType="LMT",
+         lmtPrice=msg[5]+0.25, action="SELL",
+       IBrokers::placeOrder(connib, contractobj, sell_order)
+       ## Copy new trade orders
+       eW$assign.Data("buy_id", buy_id)
+       eW$assign.Data("sell_id", sell_id)
+       ## Trade finished
+       return(eW)
```

draft: Advanced Order Types and Execution Algos

Interactive Brokers provides many different **Advanced Order Types and Execution Algos**.

The function `reqRealTimeBars()` creates data objects and then calls `tw$CALLBACK()` and passes the objects to `tw$CALLBACK()`.

`tw$CALLBACK()` calls `processMsg()` in a loop, and passes the `eWrapper()` to it.

`processMsg()` performs if-else statements and calls `eWrapper` methods.

`eWrapper()` creates an environment for data, and methods (functions) for handling (formatting) the data.

`eWrapper.RealTimeBars.CSV()` calls `eWrapper()` to create a closure (function) and returns it.

Closures allow creating mutable states for persistent data objects.

Limit trade orders can be placed inside a modified `eWrapper()` function.

```
> eWrapper_realtimebars <- function(n = 1) {
+   ew <- eWrapper_new(NULL)
+   ## ew <- IBrokers::eWrapper(NULL)
+   ew$assign.Data("data", rep(list(structure(.xts(matrix(rep(NA, rep(1, n)), 1, 1),
+   ew$realtimebars <- function(curMsg, msg, timestamp, file, ...) {
+     id <- as.numeric(msg[2])
+     file <- file[[id]]
+     data <- ew$get.Data("data")
+     attr(data[[id]], "index") <- as.numeric(msg[3])
+     nr.data <- NROW(data[[id]])
+     ## Write to file
+     cat(paste(msg[3], msg[4], msg[5], msg[6], msg[7], msg[8], msg[9], msg[10],
+     ## Write to console
+     ## ew$countn <- ew$countn + 1
+     ew$assign.Data("countn", ew$get.Data("countn")+1)
+     cat(paste0("countn=", ew$get.Data("countn"), "\tOpen=", msg[4],
+     ## cat(paste0("Open=", msg[4], "\tHigh=", msg[5], "\tLow=", msg[6],
+     ### Trade
+     ## Cancel previous trade orders
+     buy_id <- ew$get.Data("buy_id")
+     sell_id <- ew$get.Data("sell_id")
+     if (buy_id>0) IBrokers::cancelOrder(connib, buy_id)
+     if (sell_id>0) IBrokers::cancelOrder(connib, sell_id)
+     ## Execute buy limit order
+     buy_id <- IBrokers::reqIds(connib)
+     buy_order <- IBrokers::tw$Order(buy_id, orderType="LMT",
+                                   lmtPrice=msg[6]-0.25, action="BUY",
+     IBrokers::placeOrder(connib, contractobj, buy_order)
+     ## Execute sell limit order
+     sell_id <- IBrokers::reqIds(connib)
+     sell_order <- IBrokers::tw$Order(sell_id, orderType="LMT",
+                                   lmtPrice=msg[5]+0.25, action="SELL",
+     IBrokers::placeOrder(connib, contractobj, sell_order)
+     ## Copy new trade orders
+     ew$assign.Data("buy_id", buy_id)
+     ew$assign.Data("sell_id", sell_id)
+     ### Trade finished
+     return(ew$get.Data("data"))
+   }
+ }
```

draft: Downloading Live *TAQ* Data and Replaying It

Note: The script downloads the raw data, but doesn't replay the bar data properly.

<https://offerword.wordpress.com/2015/05/21/market-data-recording-and-playback-with-ibrokers-and-r-2/>

The function `reqMktData()` downloads live (real-time) trades and quotes (*TAQ*) data from Interactive Brokers.

The function `eWrapper()` formats the real-time market events (trades and quotes), so they can be displayed or saved to a file.

The method `eWrapper.MktData.CSV()` formats the real-time *TAQ* data so it can be saved to a .csv file.

The real-time *TAQ* data fields are *BidSize*, *BidPrice*, *AskPrice*, *AskSize*, *Last*, *LastSize*, *Volume*.

BidPrice is the quoted bid price, *AskPrice* is the quoted offer price, and *Last* is the most recent traded price.

The *TAQ* data is spaced irregularly in time, with data recorded each time a new trade or quote arrives.

```
> # Define S&P Emini futures June 2019 contract
> snp_contract <- IBrokers::twFuture(symbol="ES",
+   exch="GLOBEX", expiry="201906")
> # Define VIX futures June 2019 contract
> vix_contract <- IBrokers::twFuture(symbol="VIX",
+   local="VXZ8", exch="CFE", expiry="201906")
> # Define 10yr Treasury futures June 2019 contract
> trs_contract <- IBrokers::twFuture(symbol="ZN",
+   exch="ECBOT", expiry="201906")
> # Define Emini gold futures June 2019 contract
> gold_contract <- IBrokers::twFuture(symbol="YG",
+   exch="NYSELIFFE", expiry="201906")
> # Define euro currency future June 2019 contract
> euro_contract <- IBrokers::twFuture(symbol="EUR",
+   exch="GLOBEX", expiry="201906")
> IBrokers::reqContractDetails(conn=connib, Contract=euro_contract)
>
> # Define data directory
> dirn <- "/Users/jerzy/Develop/data/vix"
> # Dir.create(dirn)
>
> # Open file for error messages
> file_root <- "replay"
> filen <- file.path(dirn, paste0(file_root, "_error.csv"))
> error_connect <- file(filen, open="w")
>
> # Open file for raw data
> filen <- file.path(dirn, paste0(file_root, "_raw.csv"))
> raw_connect <- file(filen, open="w")
>
> # Create empty eWrapper to redirect error messages to error file
> error_ewrapper <- eWrapper(debug=NULL, errfile=error_connect)
>
> # Create eWrapper for raw data
> raw_ewrapper <- eWrapper(debug=TRUE)
>
> # Redirect error messages to error_ewrapper (error_ewrapper),
> # by replacing handler errorMessages() in raw_ewrapper
```

draft: Defining Instrument Pairs in *TWS*

Users can **Define Pairs of Instruments** in the *TWS*, using the *Combo Selection* window.

To open the *Combo Selection* window, right-click on a blank contract field and select *Generic Combo*.

Or enter the symbol for one of the instruments and select *Combinations* followed by *Option Combos* from the drop down menu, and then click the tab *Pair or Leg-by-leg*.

The pairs can be traded as a single instrument, but the execution is not guaranteed, and the bid-ask spread may be very large.

The tab *Multiple* allows selecting a group of combination quotes on the same underlying for comparison.

First access the *Combo Selection* window, enter the symbol for one of the instruments and select *Combinations* followed by *Option Combos* from the drop down menu, and then click the tab *Pair or Leg-by-leg*.

Interactive Brokers provides more information about financial contracts here: [IB Traded Products](#)

```
> # Define AAPL stock contract (object)
> contractobj <- IBrokers::twsequity("AAPL", primary="SMART")
> # Define CHF currency contract
> contractobj <- IBrokers::twscurrency("CHF", currency="USD")
> # Define S&P Emini future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="ES",
+   exch="GLOBEX", expiry="201906")
> # Define 10yr Treasury future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="ZN",
+   exch="ECBOT", expiry="201906")
> # Define euro currency future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="EUR",
+   exch="GLOBEX", expiry="201906")
> # Define Gold future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="GC",
+   exch="NYMEX", expiry="201906")
> # Test if contract object is correct
> IBrokers::is.twsContract(contractobj)
> # Get list with instrument information
> IBrokers::reqContractDetails(conn=connib, Contract=contractobj)
> # Install the package twsInstrument
> install.packages("twsInstrument", repos="http://r-forge.r-project.org")
> # Define euro future using getContract() and Conid
> contractobj <- twsInstrument::getContract("317631411")
> # Get list with instrument information
> IBrokers::reqContractDetails(conn=connib, Contract=contractobj)
```

draft: Defining Investment Strategies in *TWS*

Users can [Define Investment Strategies](#) in the *TWS*, using the *Portfolio Builder* window.

To open the *Portfolio Builder* window, left-click on the New Window dropdown on the upper-left of *TWS*, and select *Portfolio Builder*.

Click "Create New Strategy" to open and edit Investment Rules in the sidecar populated with sample data based on the last sorting option you selected. The main Portfolio Builder reflects changes made in the sidecar in real time.

Or enter the symbol for one of the instruments and select *Combinations* followed by *Option Combos* from the drop down menu, and then click the tab *Pair or Leg-by-leg*.

The pairs can be traded as a single instrument, but the execution is not guaranteed, and the bid-ask spread may be very large.

The tab *Multiple* allows selecting a group of combination quotes on the same underlying for comparison.

First access the *Combo Selection* window, enter the symbol for one of the instruments and select *Combinations* followed by *Option Combos* from the drop down menu, and then click the tab *Pair or Leg-by-leg*.

Interactive Brokers provides more information about

```
> # Define AAPL stock contract (object)
> contractobj <- IBrokers::twsequity("AAPL", primary="SMART")
> # Define CHF currency contract
> contractobj <- IBrokers::twscurrency("CHF", currency="USD")
> # Define S&P Emini future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="ES",
+   exch="GLOBEX", expiry="201906")
> # Define 10yr Treasury future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="ZN",
+   exch="ECBOT", expiry="201906")
> # Define euro currency future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="EUR",
+   exch="GLOBEX", expiry="201906")
> # Define Gold future June 2019 contract
> contractobj <- IBrokers::twsfuture(symbol="GC",
+   exch="NYMEX", expiry="201906")
> # Test if contract object is correct
> IBrokers::is.twsContract(contractobj)
> # Get list with instrument information
> IBrokers::reqContractDetails(conn=connib, Contract=contractobj)
> # Install the package twsInstrument
> install.packages("twsInstrument", repos="http://r-forge.r-project.")
> # Define euro future using getContract() and Conid
> contractobj <- twsInstrument::getContract("317631411")
> # Get list with instrument information
> IBrokers::reqContractDetails(conn=connib, Contract=contractobj)
```

draft: Linux

Advanced Package Tool, or apt, is a free software user interface that works with core libraries to handle the installation and removal of software on Debian, Ubuntu and other Linux distributions. Debian uses the dpkg packaging system apt is a more friendly way to handle packaging than dpkg. apt-get is similar to apt. apt consists some of the most widely used features from apt-get and apt-cache leaving aside obscure and seldom used features. apt provides the most common used commands from apt-get and apt-cache.

<https://itsfoss.com/apt-command-guide/>

<https://itsfoss.com/apt-vs-apt-get-difference/>

Get Debian version and release `cat /etc/issue` `cat /etc/os-release`

Install GNOME Debian Desktop Environment

<https://www.gnome.org/>

<https://wiki.debian.org/Gnome>

<https://wiki.debian.org/DesktopEnvironment> Install



VNC server <https://medium.com/google-cloud/graphical-user-interface-gui-for-google-compute-engine-instance-78fccda09e5c> `sudo apt-get install vnc4server vncserver`

Install Windows RDP to Debian remote desktop

<http://blog.technotesdesk.com/how-to-use-rdp-from-windows-to-connect-to-debian-or-ubuntu-machine>

apt-get install xrdp Run Windows RDP to Static IP

Address <https://cloud.google.com/compute/docs/ip-addresses/reserve-static-external-ip-address>



draft: Creating Google Cloud Project

Install R

<https://www.rstudio.com/products/rstudio/download-server/> <https://cran.rstudio.com/bin/linux/debian/>



```
sudo apt-get update sudo apt-get install r-base
r-base-dev sudo apt-get install libopenblas-base sudo
apt-get install libssl-dev sudo apt-get install libgit2-dev
sudo apt-get install libssh2-1-dev sudo apt-get install
libcurl4-openssl-dev
```

```
Change R lib permission cd /usr/local/lib/R sudo
chmod o+w site-library
```

```
Type "R" In R install.packages("openssl")
install.packages("git2r") install.packages("curl")
install.packages("httr") install.packages("usethis")
install.packages("devtools")
```

```
Upgrade R sudo apt-get remove r-base cat
/etc/apt/sources.list sudo nano /etc/apt/sources.list
deb http://cran.rstudio.com/bin/linux/debian
stretch-cran35/ sudo apt-get install dirmngr gpg
-keyserver keyserver.ubuntu.com -recv-key
FCAE2A0E115C3D8A gpg -a --export
FCAE2A0E115C3D8A — sudo apt-key add - sudo
apt-get update sudo apt-get install r-base sudo apt-get
install r-base-dev
sudo apt update sudo apt upgrade
https://www.rstudio.com/products/rstudio/download-server/
```

Connecting to Interactive Brokers via the API requires being logged into the IB Trader Workstation (TWS) or



draft: Pursuing a Career in Portfolio Management

Becoming a portfolio manager (PM) is very difficult:

- Portfolio management jobs are more difficult than they appear because of survivorship bias among portfolio managers. There are very few successful portfolio managers, and most of the others get fired. But you never hear about those.
- Becoming a portfolio manager (PM) is very difficult because it requires experience and a track record.
- Funds only want to hire experienced PMs with a track record. But how to gain experience and a track record without working as a PM?
- The best way to start is by getting a job as a quant supporting portfolio managers.
- Strong programming skills will drastically improve your chances.
- The best way to get a good job is through networking (personal contacts). Personal contacts are the most valuable resource in pursuing a career in portfolio management.
- Networking techniques: Create a Github account with your projects (it's your calling card), Collaborate on open-source projects.